

# knn

April 27, 2023

## 1 ECE 285 Assignment 1: KNN

For this part of assignment, you are tasked to implement KNN algorithm and test it on the a subset of CIFAR10 dataset.

You could run the whole notebook and answer the question in the notebook.

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

```
[1]: # Import Packages
import numpy as np
import matplotlib.pyplot as plt
```

### 1.1 Prepare Dataset

Since CIFAR10 is a relative large dataset, and KNN is quite time-consuming method, we only a small sub-set of CIFAR10 for KNN part

```
[2]: from ece285.utils.data_processing import get_cifar10_data

# Use a subset of CIFAR10 for KNN assignments
dataset = get_cifar10_data(subset_train=5000, subset_val=250, subset_test=500)

print(dataset.keys())
print("Training Set Data Shape: ", dataset["x_train"].shape)
print("Training Set Label Shape: ", dataset["y_train"].shape)
```

```
dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_test', 'y_test'])
Training Set Data Shape: (5000, 3072)
Training Set Label Shape: (5000,)
```

### 1.2 Implementation (60%)

You need to implement the KNN method in `algorithms/knn.py`. You need to fill in the prediction function(since the training of KNN is just remembering the training set).

For KNN implementation, you are tasked to implement two version of it.

- Two Loop Version: use one loop to iterate through training samples and one loop to iterate through test samples

- One Loop Version: use one loop to iterate through test samples and use broadcast feature of numpy to calculate all the distance at once

Note: It is possible to build a Fully Vectorized Version without explicit for loop to calculate the distance, but you do not have to do it in this assignment. You could use the fully vectorized version to replace the loop versions as well.

For distance function, in this assignment, we use Euclidean distance between samples.

```
[3]: from ece285.algorithms import KNN

knn = KNN(num_class=10)
knn.train(
    x_train=dataset["x_train"],
    y_train=dataset["y_train"],
    k=5,
)
```

### 1.2.1 Compare the time consumption of different method

In this section, you will test your different implementation of KNN method, and compare their speed.

```
[4]: from ece285.utils.evaluation import get_classification_accuracy
```

#### Two Loop Version:

```
[5]: import time

c_t = time.time()
prediction = knn.predict(dataset["x_test"], loop_count=2)
print("Two Loop Prediction Time:", time.time() - c_t)

test_acc = get_classification_accuracy(prediction, dataset["y_test"])
print("Test Accuracy:", test_acc)
```

Two Loop Prediction Time: 51.88799428939819

Test Accuracy: 0.278

#### One Loop Version

```
[6]: import time

c_t = time.time()
prediction = knn.predict(dataset["x_test"], loop_count=1)
print("One Loop Prediction Time:", time.time() - c_t)

test_acc = get_classification_accuracy(prediction, dataset["y_test"])
print("Test Accuracy:", test_acc)
```

One Loop Prediction Time: 57.39563202857971

Test Accuracy: 0.278

**Your different implementation should output the exact same result**

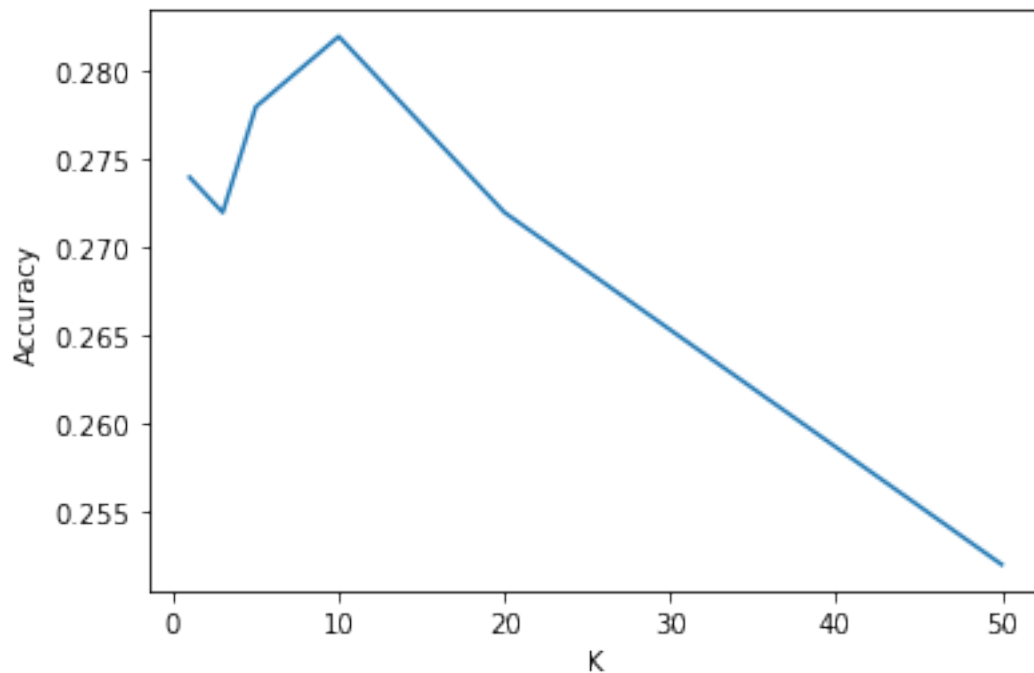
### 1.3 Test different Hyper-parameter(20%)

For KNN, there is only one hyper-parameter of the algorithm: How many nearest neighbour to use(**K**).

Here, you are provided the code to test different k for the same dataset.

```
[7]: accuracies = []

k_candidates = [1, 3, 5, 10, 20, 50]
for k_cand in k_candidates:
    prediction = knn.predict(x_test=dataset["x_test"], k=k_cand)
    acc = get_classification_accuracy(prediction, dataset["y_test"])
    accuracies.append(acc)
plt.ylabel("Accuracy")
plt.xlabel("K")
plt.plot(k_candidates, accuracies)
plt.show()
```



#### 1.3.1 Inline Question 1:

Please describe the output result you get, and provide some explanation as well.

### 1.3.2 Your Answer:

The accuracy is highest for 10 near neighbours. When  $k$  is too low, the model may overfit the training data and be sensitive to noise in the data. On the other hand, when  $k$  is too high, the model may underfit the training data and not capture the underlying patterns in the data. So in our case  $k = 10$  is a better estimate of the model.

### 1.4 Try different feature representation(20%)

Since machine learning method rely heavily on the feature extraction, you will see how different feature representation affect the performance of the algorithm in this section.

You are provided the code about using **HOG** descriptor to represent samples in the notebook.

```
[8]: from ece285.utils.data_processing import get_cifar10_data
from ece285.utils.data_processing import HOG_preprocess
from functools import partial

# Delete previous dataset to save memory
del dataset
del knn

# Use a subset of CIFAR10 for KNN assignments
hog_p_func = partial(
    HOG_preprocess,
    orientations=9,
    pixels_per_cell=(4, 4),
    cells_per_block=(1, 1),
    visualize=False,
    multichannel=True,
)
dataset = get_cifar10_data(
    feature_process=hog_p_func, subset_train=5000, subset_val=250,
    ↪subset_test=500
)
```

Start Processing


Processing Time: 118.31940770149231

```
[9]: knn = KNN(num_class=10)
knn.train(
    x_train=dataset["x_train"],
    y_train=dataset["y_train"],
    k=5,
)
accuracies = []

k_candidates = [1, 3, 5, 10, 20, 50]
for k_cand in k_candidates:
```

```
prediction = knn.predict(x_test=dataset["x_test"], k=k_cand)
acc = get_classification_accuracy(prediction, dataset["y_test"])
accuracies.append(acc)
```

```
plt.ylabel("Accuracy")
plt.xlabel("K")
plt.plot(k_candidates, accuracies)
plt.show()
```



#### 1.4.1 Inline Question 2:

Please describe the output result you get, compare with the result you get in the previous section, and provide some explanation as well.

#### 1.4.2 Your Answer:

The accuracy is higher when the value of  $k$  is 1, and then it has a sharp decrease at 3. The accuracy increases a little as the  $k$  value increases to 10, and thereafter has a gradual decrease. But on a whole, as compared to the previous model, the model using HOG descriptor to represent samples does a better performance in predicting classes. HOG uses a feature extraction and possibly reduces the redundant features in the data, thereby leading to better prediction.

# ECE 285 Assignment 1: Linear Regression

For this part of assignment, you are tasked to implement a linear regression algorithm for multiclass classification and test it on the CIFAR10 dataset.

You should run the whole notebook and answer the questions in the notebook.

CIFAR 10 dataset contains  $32 \times 32 \times 3$  RGB images of 10 distinct categories, and our aim is to predict which class the image belongs to

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

In [14]:

```
# Prepare Packages
import numpy as np
import matplotlib.pyplot as plt

from ece285.utils.data_processing import get_cifar10_data

# Use a subset of CIFAR10 for the assignment
dataset = get_cifar10_data(
    subset_train=5000,
    subset_val=250,
    subset_test=500,
)

print(dataset.keys())
print("Training Set Data Shape: ", dataset["x_train"].shape)
print("Training Set Label Shape: ", dataset["y_train"].shape)
print("Validation Set Data Shape: ", dataset["x_val"].shape)
print("Validation Set Label Shape: ", dataset["y_val"].shape)
print("Test Set Data Shape: ", dataset["x_test"].shape)
print("Test Set Label Shape: ", dataset["y_test"].shape)
```

```
dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_
test', 'y_test'])
Training Set Data Shape: (5000, 3072)
Training Set Label Shape: (5000,)
Validation Set Data Shape: (250, 3072)
Validation Set Label Shape: (250,)
Test Set Data Shape: (500, 3072)
Test Set Label Shape: (500,)
```

In [15]:

```
x_train = dataset["x_train"]  
y_train = dataset["y_train"]  
x_val = dataset["x_val"]  
y_val = dataset["y_val"]  
x_test = dataset["x_test"]  
y_test = dataset["y_test"]
```



In [16]:

```
# Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = [
    "plane",
    "car",
    "bird",
    "cat",
    "deer",
    "dog",
    "frog",
    "horse",
    "ship",
    "truck",
]

samples_per_class = 7

def visualize_data(dataset, classes, samples_per_class):
    num_classes = len(classes)
    for y, cls in enumerate(classes):
        idxs = np.flatnonzero(y_train == y)
        idxs = np.random.choice(idxs, samples_per_class, replace=False)
        for i, idx in enumerate(idxs):
            plt_idx = i * num_classes + y + 1
            plt.subplot(samples_per_class, num_classes, plt_idx)
            plt.imshow(dataset[idx])
            plt.axis("off")
            if i == 0:
                plt.title(cls)
    plt.show()

visualize_data(
    x_train.reshape(5000, 3, 32, 32).transpose(0, 2, 3, 1), classes
)
```



# Linear Regression for multi-class classification

A Linear Regression Algorithm has 2 hyperparameters that you can experiment with:

- **Learning rate** - controls how much we change the current weights of the classifier during each update. We set it at a default value of 0.5, and later you are asked to experiment with different values. We recommend looking at the graphs and observing how the performance of the classifier changes with different learning rate.
- **Number of Epochs** - An epoch is a complete iterative pass over all of the data in the dataset. During an epoch we predict a label using the classifier and then update the weights of the classifier according the linear classifier update rule for each sample in the training set. We evaluate our models after every 10 epochs and save the accuracies, which are later used to plot the training, validation and test VS epoch curves.
- **Weight Decay** - Regularization can be used to constrain the weights of the classifier and prevent their values from blowing up. Regularization helps in combatting overfitting. You will be using the 'weight\_decay' term to introduce regularization in the classifier.

## Implementation (50%)

You first need to implement the Linear Regression method in

`algorithms/linear_regression.py`. You need to fill in the training function as well as the prediction function.

In [17]:

```
# Import the algorithm implementation (TODO: Complete the Linear Re
from ece285.algorithms import Linear
from ece285.utils.evaluation import get_classification_accuracy

num_classes = 10 # Cifar10 dataset has 10 different classes

# Initialize hyper-parameters
learning_rate = 0.0001 # You will be later asked to experiment wit
num_epochs_total = 1000 # Total number of epochs to train the clas
epochs_per_evaluation = 10 # Epochs per step of evaluation; We wil
N, D = dataset[
    "x_train"
].shape # Get training data shape, N: Number of examples, D:Dimens
weight_decay = 0.0

# Insert additional scalar term 1 in the samples to account for the
x_train = np.insert(x_train, D, values=1, axis=1)
x_val = np.insert(x_val, D, values=1, axis=1)
x_test = np.insert(x_test, D, values=1, axis=1)
```

In [18]:

```
# Training and evaluation function -> Outputs accuracy data
def train(learning_rate_, weight_decay_):
    # Create a linear regression object
    linear_regression = Linear(
        num_classes, learning_rate_, epochs_per_evaluation, weight_
    )

    # Randomly initialize the weights and biases
    weights = np.random.randn(num_classes, D + 1) * 0.0001

    train_accuracies, val_accuracies, test_accuracies = [], [], []

    # Train the classifier
    for _ in range(int(num_epochs_total / epochs_per_evaluation)):
        # Train the classifier on the training data
        weights = linear_regression.train(x_train, y_train, weights)

        # Evaluate the trained classifier on the training dataset
        y_pred_train = linear_regression.predict(x_train)
        train_accuracies.append(get_classification_accuracy(y_pred_train, y_train))

        # Evaluate the trained classifier on the validation dataset
        y_pred_val = linear_regression.predict(x_val)
        val_accuracies.append(get_classification_accuracy(y_pred_val, y_val))

        # Evaluate the trained classifier on the test dataset
        y_pred_test = linear_regression.predict(x_test)
        test_accuracies.append(get_classification_accuracy(y_pred_test, y_test))

    return train_accuracies, val_accuracies, test_accuracies, weights
```

## Plot the Accuracies vs epoch graphs

In [19]:

```
import matplotlib.pyplot as plt


def plot_accuracies(train_acc, val_acc, test_acc):
    # Plot Accuracies vs Epochs graph for all the three
    epochs = np.arange(0, int(num_epochs_total / epochs_per_evaluation))
    plt.ylabel("Accuracy")
    plt.xlabel("Epoch/10")
    plt.plot(epochs, train_acc, epochs, val_acc, epochs, test_acc)
    plt.legend(["Training", "Validation", "Testing"])
    plt.show()
```

In [20]:

```
# Run training and plotting for default parameter values as mentioned
t_ac, v_ac, te_ac, weights = train(learning_rate, weight_decay)
```

In [21]:

```
plot_accuracies(t_ac, v_ac, te_ac)
```



## Try different learning rates and plot graphs for all (20%)

In [22]:


```
# Initialize the best values
best_weights = weights
best_learning_rate = learning_rate
best_weight_decay = weight_decay



# TODO
# Repeat the above training and evaluation steps for the following
# You need to try 3 learning rates and submit all 3 graphs along wi
learning_rates = [0.00005, 0.0001, 0.0005]
weight_decay = 0.0 # No regularization for now

# FEEL FREE TO EXPERIMENT WITH OTHER VALUES. REPORT OTHER VALUES IF

# for lr in learning_rates: Train the classifier and plot data
# Step 1. train_accu, val_accu, test_accu = train(lr, weight_decay)
# Step 2. plot accuracies(train_accu, val_accu, test_accu)

for learning_rate in learning_rates:
    # TODO: Train the classifier with different learning rates and
    t_ac, v_ac, te_ac, weights = train(learning_rate, weight_decay)
    plot accuracies(t_ac, v_ac, te_ac)
```





### Inline Question 1.

Which one of these learning rates (best\_lr) would you pick to train your model? Please Explain why.

**Your Answer:**



The best learning rate out of the above code is 0.0005, and hence that would be a better pick to train the model. If the learning rate is too small, the model might converge very slowly, leading to a longer training time and less efficient optimization. Conversely, if the learning rate is too large, the algorithm might overshoot the best solution and diverge.

## Regularization: Try different weight decay and plot graphs for all (20%)

In [23]:


```
# Initialize a non-zero weight_decay (Regularization constant) term
# Use the best learning rate as obtained from the above exercise, b

# You need to try 3 learning rates and submit all 3 graphs along wi
weight_decays = [0.001,0.01,0.1]

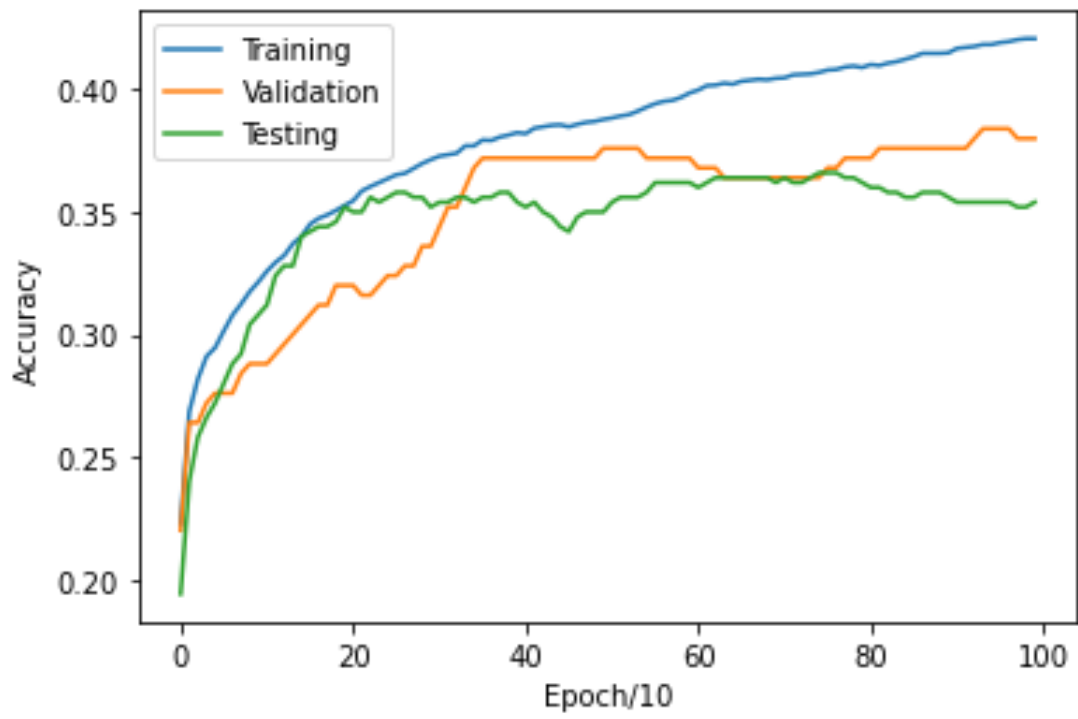
# FEEL FREE TO EXPERIMENT WITH OTHER VALUES. REPORT OTHER VALUES IF

# for weight_decay in weight_decays: Train the classifier and plot
# Step 1. train_accu, val_accu, test_accu = train(best_lr, weight_d
# Step 2. plot accuracies(train_accu, val_accu, test_accu)


for weight_decay in weight_decays:
    # TODO: Train the classifier with different weighty decay and p
    t_ac, v_ac, te_ac, weights = train(learning_rate, weight_decay)
    plot accuracies(t_ac, v_ac, te_ac)
    print(t_ac[-1], v_ac[-1], te_ac[-1])
```



0.421 0.38 0.352



0.4208 0.38 0.354



0.4204 0.38 0.354

**Inline Question 2.**

### **Your Answer:**

Underfitting occurs when the model is too simple, while overfitting occurs when it is too complex. The `weight_decay` term of 0.1 gives the best classifier performance as it balances the trade-off between underfitting and overfitting, leading to good generalization performance on testing data.

**Visualize the filters (10%)**

In [25]:

```
# These visualizations will only somewhat make sense if your learning rate is
# properly chosen in the model. Do your best.

# TODO: Run this cell and Show filter visualizations for the best set of hyperparameters
# Report the 2 hyperparameters you used to obtain the best model.
best_learning_rate = 0.0005
best_weight_decay = 0.1
t_ac, v_ac, te_ac, best_weights = train(best_learning_rate, best_weight_decay)

# NOTE: You need to set `best_learning_rate` and `best_weight_decay` to the values
print("Best LR:", best_learning_rate)
print("Best Weight Decay:", best_weight_decay)

# NOTE: You need to set `best_weights` to the weights with the highest validation accuracy
w = best_weights[:, :-1]
w = w.reshape(10, 3, 32, 32).transpose(0, 2, 3, 1)

w_min, w_max = np.min(w), np.max(w)

fig = plt.figure(figsize=(20, 20))
classes = [
    "plane",
    "car",
    "bird",
    "cat",
    "deer",
    "dog",
    "frog",
    "horse",
    "ship",
    "truck",
]

for i in range(10):
    fig.add_subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[i, :, :, :].squeeze() - w_min) / (w_max - w_min)
    # plt.imshow(wimg.astype('uint8'))
    plt.imshow(wimg.astype(int))
    plt.axis("off")
    plt.title(classes[i])

plt.show()
```

Best LR: 0.0005

Best Weight Decay: 0.1


plane



car



bird




cat




deer



dog




frog



horse



ship



truck



# logistic\_regression

April 28, 2023

## 1 ECE 285 Assignment 1: Logistic Regression

For this part of assignment, you are tasked to implement a logistic regression algorithm for multi-class classification and test it on the CIFAR10 dataset.

You could run the whole notebook and answer the questions in the notebook.

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

```
[1]: # Prepare Packages
import numpy as np
import matplotlib.pyplot as plt

from ece285.utils.data_processing import get_cifar10_data

# Use a subset of CIFAR10 for KNN assignments
dataset = get_cifar10_data(
    subset_train=5000,
    subset_val=250,
    subset_test=500,
)

print(dataset.keys())
print("Training Set Data Shape: ", dataset["x_train"].shape)
print("Training Set Label Shape: ", dataset["y_train"].shape)
print("Validation Set Data Shape: ", dataset["x_val"].shape)
print("Validation Set Label Shape: ", dataset["y_val"].shape)
print("Test Set Data Shape: ", dataset["x_test"].shape)
print("Test Set Label Shape: ", dataset["y_test"].shape)
```

```
dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_test', 'y_test'])
Training Set Data Shape: (5000, 3072)
Training Set Label Shape: (5000,)
Validation Set Data Shape: (250, 3072)
Validation Set Label Shape: (250,)
Test Set Data Shape: (500, 3072)
Test Set Label Shape: (500,)
```



## 2 Logistic Regression for multi-class classification

A Logistic Regression Algorithm has 3 hyperparameters that you can experiment with:

- **Learning rate** - controls how much we change the current weights of the classifier during each update. We set it at a default value of 0.5, and later you are asked to experiment with different values. We recommend looking at the graphs and observing how the performance of the classifier changes with different learning rate.
- **Number of Epochs** - An epoch is a complete iterative pass over all of the data in the dataset. During an epoch we predict a label using the classifier and then update the weights of the classifier according the linear classifier update rule for each sample in the training set. We evaluate our models after every 10 epochs and save the accuracies, which are later used to plot the training, validation and test VS epoch curves.
- **Weight Decay** - Regularization can be used to constrain the weights of the classifier and prevent their values from blowing up. Regularization helps in combatting overfitting. You will be using the 'weight\_decay' term to introduce regularization in the classifier.

The only way how a Logistic Regression based classification algorithm is different from a Linear Regression algorithm is that in the former we additionally pass the classifier outputs into a sigmoid function which squashes the output in the (0,1) range. Essentially these values then represent the probabilities of that sample belonging to class particular classes

### 2.0.1 Implementation (40%)

You need to implement the Linear Regression method in `algorithms/logistic_regression.py`. You need to fill in the sigmoid function, training function as well as the prediction function.

```
[2]: # Import the algorithm implementation (TODO: Complete the Logistic Regression
      ↪ in algorithms/logistic_regression.py)
from ece285.algorithms import Logistic
from ece285.utils.evaluation import get_classification_accuracy

num_classes = 10 # Cifar10 dataset has 10 different classes

# Initialize hyper-parameters
learning_rate = 0.01 # You will be later asked to experiment with different
                    ↪ learning rates and report results
num_epochs_total = 1000 # Total number of epochs to train the classifier
epochs_per_evaluation = 10 # Epochs per step of evaluation; We will evaluate
                        ↪ our model regularly during training
N, D = dataset[
    "x_train"
].shape # Get training data shape, N: Number of examples, D: Dimensionality of
        ↪ the data
weight_decay = 0.00002

x_train = dataset["x_train"].copy()
y_train = dataset["y_train"].copy()
x_val = dataset["x_val"].copy()
```

```

y_val = dataset["y_val"].copy()
x_test = dataset["x_test"].copy()
y_test = dataset["y_test"].copy()

# Insert additional scalar term 1 in the samples to account for the bias as
↳discussed in class
x_train = np.insert(x_train, D, values=1, axis=1)
x_val = np.insert(x_val, D, values=1, axis=1)
x_test = np.insert(x_test, D, values=1, axis=1)

```

```

[3]: # Training and evaluation function -> Outputs accuracy data
def train(learning_rate_, weight_decay_):
    # Create a linear regression object
    logistic_regression = Logistic(
        num_classes, learning_rate_, epochs_per_evaluation, weight_decay_
    )

    # Randomly initialize the weights and biases
    weights = np.random.randn(num_classes, D + 1) * 0.0001

    train_accuracies, val_accuracies, test_accuracies = [], [], []

    # Train the classifier
    for _ in range(int(num_epochs_total / epochs_per_evaluation)):
        # Train the classifier on the training data
        weights = logistic_regression.train(x_train, y_train, weights)

        # Evaluate the trained classifier on the training dataset
        y_pred_train = logistic_regression.predict(x_train)
        train_accuracies.append(get_classification_accuracy(y_pred_train,
↳y_train))
        #print(get_classification_accuracy(y_pred_train, y_train))

        # Evaluate the trained classifier on the validation dataset
        y_pred_val = logistic_regression.predict(x_val)
        val_accuracies.append(get_classification_accuracy(y_pred_val, y_val))

        # Evaluate the trained classifier on the test dataset
        y_pred_test = logistic_regression.predict(x_test)
        test_accuracies.append(get_classification_accuracy(y_pred_test, y_test))

    return train_accuracies, val_accuracies, test_accuracies, weights

```

```

[4]: import matplotlib.pyplot as plt

def plot_accuracies(train_acc, val_acc, test_acc):

```

```

# Plot Accuracies vs Epochs graph for all the three
epochs = np.arange(0, int(num_epochs_total / epochs_per_evaluation))
plt.ylabel("Accuracy")
plt.xlabel("Epoch/10")
plt.plot(epochs, train_acc, epochs, val_acc, epochs, test_acc)
plt.legend(["Training", "Validation", "Testing"])
plt.show()

```

```


[ ]: # Run training and plotting for default parameter values as mentioned above
t_ac, v_ac, te_ac, weights = train(learning_rate, weight_decay)

```

```

[ ]: plot_accuracies(t_ac, v_ac, te_ac)

```



## 2.0.2 Try different learning rates and plot graphs for all (20%)

```

[7]: # Initialize the best values
best_weights = weights
best_learning_rate = learning_rate
best_weight_decay = weight_decay

# TODO
# Repeat the above training and evaluation steps for the following learning
→ rates and plot graphs

```

```


# You need to try 3 learning rates and submit all 3 graphs along with this_
↪notebook pdf to show your learning rate experiments
learning_rates = [0.005, 0.01, 0.1]
weight_decay = 0.0 # No regularization for now

# FEEL FREE TO EXPERIMENT WITH OTHER VALUES. REPORT OTHER VALUES IF THEY_
↪ACHIEVE A BETTER PERFORMANCE


# for lr in learning_rates: Train the classifier and plot data
# Step 1. train_accu, val_accu, test_accu = train(lr, weight_decay)
# Step 2. plot accuracies(train_accu, val_accu, test_accu)

for learning_rate in learning_rates:
    # TODO: Train the classifier with different learning rates and plot
    t_ac, v_ac, te_ac, weights = train(learning_rate, weight_decay)
    plot accuracies(t_ac, v_ac, te_ac)
    print(t_ac[-1], v_ac[-1], te_ac[-1])


```



0.397 0.368 0.332



0.4204 0.396 0.336



0.1344 0.108 0.12

**Inline Question 1.** Which one of these learning rates (best\_lr) would you pick to train your model? Please Explain why.

**Your Answer:** The best learning rate out of the above code is 0.01, and hence that would be a better pick to train the model. If the learning rate is too small, the model might converge very slowly, leading to a longer training time and less efficient optimization. Conversely, if the learning rate is too large, the algorithm might overshoot the best solution and diverge, causing the model to perform poorly. In our case out of the models tested 0.01 seems to be a better fit.

### 2.0.3 Regularization: Try different weight decay and plots graphs for all (20%)


```
[8]: # Initialize a non-zero weight_decay (Regularization constant) term and repeat
    ↳ the training and evaluation
    # Use the best learning rate as obtained from the above exercise, best_lr

    # You need to try 3 learning rates and submit all 3 graphs along with this
    ↳ notebook pdf to show your weight decay experiments
weight_decays = [0.00005, 0.00002, 0.00001]
learning_rate = 0.01


    # FEEL FREE TO EXPERIMENT WITH OTHER VALUES. REPORT OTHER VALUES IF THEY
    ↳ ACHIEVE A BETTER PERFORMANCE

    # for weight_decay in weight_decays: Train the classifier and plot data
    # Step 1. train_accu, val_accu, test_accu = train(best_lr, weight_decay)
    # Step 2. plot accuracies(train_accu, val_accu, test_accu)


for weight_decay in weight_decays:
    # TODO: Train the classifier with different weight decay and plot
    t_ac, v_ac, te_ac, weights = train(learning_rate, weight_decay)
    plot accuracies(t_ac, v_ac, te_ac)
    print(t_ac[-1], v_ac[-1], te_ac[-1])
```



0.4202 0.396 0.336



0.4204 0.396 0.336



0.4202 0.396 0.336

**Inline Question 2.** Discuss underfitting and overfitting as observed in the 3 graphs obtained by changing the regularization. Which `weight_decay` term gave you the best classifier performance? HINT: Do not just think in terms of best training set performance, keep in mind that the real utility of a machine learning model is when it performs well on data it has never seen before

**Your Answer:** Underfitting occurs when the model is too simple, while overfitting occurs when it is too complex. All the `weight_decay` terms here gives the same classifier performance as it balances the trade-off between underfitting and overfitting, leading to good generalization performance on testing data.

## 2.0.4 Visualize the filters (10%)

```
[9]: # These visualizations will only somewhat make sense if your learning rate and
      ↪ weight_decay parameters were
      # properly chosen in the model. Do your best.

      # TODO: Run this cell and Show filter visualizations for the best set of
      ↪ weights you obtain.
      # Report the 2 hyperparameters you used to obtain the best model.

      best_learning_rate = 0.001
      best_weight_decay = 0.00002
      t_ac, v_ac, te_ac, best_weights = train(learning_rate, weight_decay)
```



```

# NOTE: You need to set `best_learning_rate` and `best_weight_decay` to the
↳ values that gave the highest accuracy
print("Best LR:", best_learning_rate)
print("Best Weight Decay:", best_weight_decay)

# NOTE: You need to set `best_weights` to the weights with the highest accuracy
w = best_weights[:, :-1]
w = w.reshape(10, 3, 32, 32).transpose(0, 2, 3, 1)

w_min, w_max = np.min(w), np.max(w)


fig = plt.figure(figsize=(16, 16))
classes = [
    "plane",
    "car",
    "bird",
    "cat",
    "deer",
    "dog",
    "frog",
    "horse",
    "ship",
    "truck",
]
for i in range(10):
    fig.add_subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[i, :, :, :].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype(int))
    plt.axis("off")
    plt.title(classes[i])
plt.show()

```

Best LR: 0.001

Best Weight Decay: 2e-05



**Inline Question 3. (10%)**

- Compare and contrast the performance of the 2 classifiers i.e. Linear Regression and Logistic Regression.
- Which classifier would you deploy for your multiclass classification project and why?

**Your Answer:** Logistic regression performs better than the linear regression. For classification, mostly always Logistic regression better performs the linear regression. For working on a multiclass classification project, logistic regression is likely the better choice due to its ability to output probability scores and handle non-linear relationships.

# neural\_network

April 28, 2023

## 1 ECE285 Assignment 1: Neural Network in NumPy

Use this notebook to build your neural network by implementing the following functions in the python files under `ece285/algorithms` directory:

1. `linear.py`
2. `relu.py`
3. `softmax.py`
4. `loss_func.py`

You will be testing your 2 layer neural network implementation on a toy dataset.

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

```
[1]: # Setup
import matplotlib.pyplot as plt
import numpy as np

from ece285.layers.sequential import Sequential
from ece285.layers.linear import Linear
from ece285.layers.relu import ReLU
from ece285.layers.softmax import Softmax
from ece285.layers.loss_func import CrossEntropyLoss
from ece285.utils.optimizer import SGD

%matplotlib inline
plt.rcParams["figure.figsize"] = (10.0, 8.0) # set default size of plots

# For auto-reloading external modules
# See http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

We will use the class `Sequential` as implemented in the file `assignment2/layers/sequential.py` to build a layer by layer model of our neural network. Below we initialize the toy model and the toy random data that you will use to develop your implementation.

```
[2]: # Create a small net and some toy data to check your implementations.  
# Note that we set the random seed for repeatable experiments.
```

```
input_size = 4  
hidden_size = 10  
num_classes = 3 # Output  
num_inputs = 10 # N  
  
def init_toy_model():  
    np.random.seed(0)  
    l1 = Linear(input_size, hidden_size)  
    l2 = Linear(hidden_size, num_classes)  
  
    r1 = ReLU()  
    softmax = Softmax()  
    return Sequential([l1, r1, l2, softmax])  
  
def init_toy_data():  
    np.random.seed(0)  
    X = 10 * np.random.randn(num_inputs, input_size)  
    y = np.random.randint(num_classes, size=num_inputs)  
    # y = np.array([0, 1, 2, 2, 1])  
    return X, y  
  
net = init_toy_model()  
X, y = init_toy_data()
```

### 1.0.1 Forward Pass: Compute Scores (20%)

Implement the forward functions in Linear, Relu and Softmax layers and get the output by passing our toy data X. The output must match the given output scores

```
[3]: scores = net.forward(X)  
print("Your scores:")  
print(scores)  
print()  
print("correct scores:")  
correct_scores = np.asarray(  
    [  
        [0.33333514, 0.33333826, 0.33332661],  
        [0.3333351, 0.33333828, 0.33332661],  
        [0.3333351, 0.33333828, 0.33332662],  
        [0.3333351, 0.33333828, 0.33332662],  
        [0.33333509, 0.33333829, 0.33332662],
```

```

        [0.33333508, 0.33333829, 0.33332662],
        [0.33333511, 0.33333828, 0.33332661],
        [0.33333512, 0.33333827, 0.33332661],
        [0.33333508, 0.33333829, 0.33332662],
        [0.33333511, 0.33333828, 0.33332662],
    ]
)
print(correct_scores)

# The difference should be very small. We get < 1e-7
print("Difference between your scores and correct scores:")
print(np.sum(np.abs(scores - correct_scores)))

```

Your scores:

```

[[0.33333514 0.33333826 0.33332661]
 [0.3333351  0.33333828 0.33332661]
 [0.3333351  0.33333828 0.33332662]
 [0.3333351  0.33333828 0.33332662]
 [0.33333509 0.33333829 0.33332662]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332661]
 [0.33333512 0.33333827 0.33332661]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332662]]

```

correct scores:

```

[[0.33333514 0.33333826 0.33332661]
 [0.3333351  0.33333828 0.33332661]
 [0.3333351  0.33333828 0.33332662]
 [0.3333351  0.33333828 0.33332662]
 [0.33333509 0.33333829 0.33332662]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332661]
 [0.33333512 0.33333827 0.33332661]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332662]]

```

Difference between your scores and correct scores:

8.799388540037256e-08

## 1.0.2 Forward Pass: Compute loss given the output scores from the previous step (10%)

Implement the forward function in the `loss_func.py` file, and output the loss value. The loss value must match the given loss value.

```

[4]: Loss = CrossEntropyLoss()
     loss = Loss.forward(scores, y)
     correct_loss = 1.098612723362578

```

```

print(loss)
# should be very small, we get < 1e-12
print("Difference between your loss and correct loss:")
print(np.sum(np.abs(loss - correct_loss)))

```

```

1.0986127233625778
Difference between your loss and correct loss:
2.220446049250313e-16

```

```
[5]: y
```

```
[5]: array([2, 1, 0, 1, 2, 0, 0, 2, 0, 0])
```

### 1.0.3 Backward Pass (40%)

Implement the rest of the functions in the given files. Specifically, implement the backward function in all the 4 files as mentioned in the files. Note: No backward function in the softmax file, the gradient for softmax is jointly calculated with the cross entropy loss in the `loss_func.backward` function.

You will use the chain rule to calculate gradient individually for each layer. You can assume that this calculated gradient then is passed to the next layers in a reversed manner due to the Sequential implementation. So all you need to worry about is implementing the gradient for the current layer and multiply it with the incoming gradient (passed to the backward function as `dout`) to calculate the total gradient for the parameters of that layer.

We check the values for these gradients by calculating the difference, it is expected to get difference  $< 1e-8$ .

```

[6]: # No need to edit anything in this block ( 20% of the above 40% )
net.backward(Loss.backward())

gradients = []
for module in net._modules:
    for para, grad in zip(module.parameters, module.grads):
        assert grad is not None, "No Gradient"
        # Print gradients of the linear layer
        print(grad.shape)
        gradients.append(grad)

# Check shapes of your gradient. Note that only the linear layer has parameters
# (4, 10) -> Layer 1 W
# (10,)    -> Layer 1 b
# (10, 3)  -> Layer 2 W
# (3,)     -> Layer 2 b

```

```

(4, 10)
(10,)

```

(10, 3)  
(3,)

[7]: *# No need to edit anything in this block ( 20% of the above 40% )*

```
grad_w1 = np.array(  
    [  
        [  
            -6.24320917e-05,  
            3.41037180e-06,  
            -1.69125969e-05,  
            2.41514079e-05,  
            3.88697976e-06,  
            7.63842314e-05,  
            -8.88925758e-05,  
            3.34909890e-05,  
            -1.42758303e-05,  
            -4.74748560e-06,  
        ],  
        [  
            -7.16182867e-05,  
            4.63270039e-06,  
            -2.20344270e-05,  
            -2.72027034e-06,  
            6.52903437e-07,  
            8.97294847e-05,  
            -1.05981609e-04,  
            4.15825391e-05,  
            -2.12210745e-05,  
            3.06061658e-05,  
        ],  
        [  
            -1.69074923e-05,  
            -8.83185056e-06,  
            3.10730840e-05,  
            1.23010428e-05,  
            5.25830316e-05,  
            -7.82980115e-06,  
            3.02117990e-05,  
            -3.37645284e-05,  
            6.17276346e-05,  
            -1.10735656e-05,  
        ],  
        [  
            -4.35902272e-05,  
            3.71512704e-06,  
            -1.66837877e-05,  
            2.54069557e-06,  
        ]  
    ]  
)
```

```

        -4.33258099e-06,
        5.72310022e-05,
        -6.94881762e-05,
        2.92408329e-05,
        -1.89369767e-05,
        2.01692516e-05,
    ],
]
)
grad_b1 = np.array(
    [
        -2.27150209e-06,
        5.14674340e-07,
        -2.04284403e-06,
        6.08849787e-07,
        -1.92177796e-06,
        3.92085824e-06,
        -5.40772636e-06,
        2.93354593e-06,
        -3.14568138e-06,
        5.27501592e-11,
    ]
)

grad_w2 = np.array(
    [
        [1.28932983e-04, 1.19946731e-04, -2.48879714e-04],
        [1.08784150e-04, 1.55140199e-04, -2.63924349e-04],
        [6.96017544e-05, 1.42748410e-04, -2.12350164e-04],
        [9.92512487e-05, 1.73257611e-04, -2.72508860e-04],
        [2.05484895e-05, 4.96161144e-05, -7.01646039e-05],
        [8.20539510e-05, 9.37063861e-05, -1.75760337e-04],
        [2.45831715e-05, 8.74369112e-05, -1.12020083e-04],
        [1.34073379e-04, 1.86253064e-04, -3.20326443e-04],
        [8.86473128e-05, 2.35554414e-04, -3.24201726e-04],
        [3.57433149e-05, 1.91164061e-04, -2.26907376e-04],
    ]
)

grad_b2 = np.array([-0.1666649, 0.13333828, 0.03332662])

difference = (
    np.sum(np.abs(gradients[0] - grad_w1))
    + np.sum(np.abs(gradients[1] - grad_b1))
    + np.sum(np.abs(gradients[2] - grad_w2))
    + np.sum(np.abs(gradients[3] - grad_b2))
)

```



```
print("Difference in Gradient values", difference)
```

Difference in Gradient values 7.701916434367482e-09

```
[8]: gradients[3].shape
```

```
[8]: (3,)
```

## 1.1 Train the complete network on the toy data. (30%)

To train the network we will use stochastic gradient descent (SGD), we have implemented the optimizer for you. You do not implement any more functions in the python files. Below we implement the training procedure, you should get yourself familiar with the training process. Specifically looking at which functions to call and when.

Once you have implemented the method and tested various parts in the above blocks, run the code below to train a two-layer network on toy data. You should see your training loss decrease below 0.01.

```
[9]: # Training Procedure
# Initialize the optimizer. DO NOT change any of the hyper-parameters here or
#   above.
# We have implemented the SGD optimizer class for you here, which visits each
#   layer sequentially to
# get the gradients and optimize the respective parameters.
# You should work with the given parameters and only edit your implementation
#   in the .py files

epochs = 1000
optim = SGD(net, lr=0.1, weight_decay=0.00001)

epoch_loss = []
for epoch in range(epochs):
    # Get output scores from the network
    output_x = net(X)
    # Calculate the loss for these output scores, given the true labels
    loss = Loss.forward(output_x, y)
    # Initialize your gradients to None in each epoch
    optim.zero_grad()
    # Make a backward pass to update the internal gradients in the layers
    net.backward(Loss.backward())
    # call the step function in the optimizer to update the values of the
    #   params with the gradients
    optim.step()
    # Append the loss at each iteration
    epoch_loss.append(loss)

    if (epoch + 1) % 50 == 0:
```

```
print("Epoch {}, loss={:3f}".format(epoch + 1, epoch_loss[-1]))
```

```
Epoch 50, loss=0.832706
Epoch 100, loss=0.454687
Epoch 150, loss=0.118350
Epoch 200, loss=0.055911
Epoch 250, loss=0.038039
Epoch 300, loss=0.029528
Epoch 350, loss=0.024400
Epoch 400, loss=0.020819
Epoch 450, loss=0.017947
Epoch 500, loss=0.015866
Epoch 550, loss=0.014198
Epoch 600, loss=0.012916
Epoch 650, loss=0.011859
Epoch 700, loss=0.010943
Epoch 750, loss=0.010198
Epoch 800, loss=0.009540
Epoch 850, loss=0.008970
Epoch 900, loss=0.008454
Epoch 950, loss=0.008003
Epoch 1000, loss=0.007593
```

```
[10]: # Test your predictions. The predictions must match the labels
print(net.predict(X))
print(y)
```


```
[2 1 0 1 2 0 0 2 0 0]
[2 1 0 1 2 0 0 2 0 0]
```

```
[11]: # You should be able to achieve a training loss of less than 0.02 (10%)
print("Final training loss", epoch_loss[-1])
```

```
Final training loss 0.00759341980173128
```

```
[12]: # Plot the training loss curve. The loss in the curve should be decreasing (20%)
plt.plot(epoch_loss)
plt.title("Loss history")
plt.xlabel("Iteration")
plt.ylabel("Loss")
```

```
[12]: Text(0, 0.5, 'Loss')
```



# classification\_nn

April 28, 2023

## 1 ECE 285 Assignment 1: Classification using Neural Network

Now that you have developed and tested your model on the toy dataset set. It's time to get down and get dirty with a standard dataset such as cifar10. At this point, you will be using the provided training data to tune the hyper-parameters of your network such that it works with cifar10 for the task of multi-class classification.

Important: Recall that now we have non-linear decision boundaries, thus we do not need to do one vs all classification. We learn a single non-linear decision boundary instead. Our non-linear boundaries (thanks to relu non-linearity) will take care of differentiating between all the classes

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

```
[1]: # Prepare Packages
import numpy as np
import matplotlib.pyplot as plt

from ece285.utils.data_processing import get_cifar10_data
from ece285.utils.evaluation import get_classification_accuracy

%matplotlib inline
plt.rcParams["figure.figsize"] = (10.0, 8.0) # set default size of plots

# For auto-reloading external modules
# See http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

# Use a subset of CIFAR10 for the assignment
dataset = get_cifar10_data(
    subset_train=5000,
    subset_val=250,
    subset_test=500,
)

print(dataset.keys())
print("Training Set Data Shape: ", dataset["x_train"].shape)
```

```

print("Training Set Label Shape: ", dataset["y_train"].shape)
print("Validation Set Data Shape: ", dataset["x_val"].shape)
print("Validation Set Label Shape: ", dataset["y_val"].shape)
print("Test Set Data Shape: ", dataset["x_test"].shape)
print("Test Set Label Shape: ", dataset["y_test"].shape)

```

```

dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_test', 'y_test'])
Training Set Data Shape: (5000, 3072)
Training Set Label Shape: (5000,)
Validation Set Data Shape: (250, 3072)
Validation Set Label Shape: (250,)
Test Set Data Shape: (500, 3072)
Test Set Label Shape: (500,)

```

```

[2]: x_train = dataset["x_train"]
     y_train = dataset["y_train"]
     x_val = dataset["x_val"]
     y_val = dataset["y_val"]
     x_test = dataset["x_test"]
     y_test = dataset["y_test"]

```

```

[3]: # Import more utilities and the layers you have implemented
     from ece285.layers.sequential import Sequential
     from ece285.layers.linear import Linear
     from ece285.layers.relu import ReLU
     from ece285.layers.softmax import Softmax
     from ece285.layers.loss_func import CrossEntropyLoss
     from ece285.utils.optimizer import SGD
     from ece285.utils.dataset import DataLoader
     from ece285.utils.trainer import Trainer

```

## 1.1 Visualize some examples from the dataset.

```

[4]: # We show a few examples of training images from each class.
     classes = [
         "airplane",
         "automobile",
         "bird",
         "cat",
         "deer",
         "dog",
         "frog",
         "horse",
         "ship",
     ]
     samples_per_class = 7

```

```

def visualize_data(dataset, classes, samples_per_class):
    num_classes = len(classes)
    for y, cls in enumerate(classes):
        idxs = np.flatnonzero(y_train == y)
        idxs = np.random.choice(idxs, samples_per_class, replace=False)
        for i, idx in enumerate(idxs):
            plt_idx = i * num_classes + y + 1
            plt.subplot(samples_per_class, num_classes, plt_idx)
            plt.imshow(dataset[idx])
            plt.axis("off")
            if i == 0:
                plt.title(cls)
    plt.show()

# Visualize the first 10 classes
visualize_data(
    x_train.reshape(5000, 3, 32, 32).transpose(0, 2, 3, 1),
    classes,
    samples_per_class,
)

```



## 1.2 Initialize the model

```
[5]: input_size = 3072
hidden_size = 100 # Hidden layer size (Hyper-parameter)
num_classes = 10 # Output

# For a default setting we use the same model we used for the toy dataset.
# This tells you the power of a 2 layered Neural Network. Recall the Universal
    ↳ Approximation Theorem.
# A 2 layer neural network with non-linearities can approximate any function,
    ↳ given large enough hidden layer
def init_model():
    # np.random.seed(0) # No need to fix the seed here
    l1 = Linear(input_size, hidden_size)
    l2 = Linear(hidden_size, num_classes)

    r1 = ReLU()
    softmax = Softmax()
```

```
return Sequential([l1, r1, l2, softmax])
```

```
[6]: # Initialize the dataset with the dataloader class
dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
net = init_model()
optim = SGD(net, lr=0.01, weight_decay=0.01)
loss_func = CrossEntropyLoss()
epoch = 200 # (Hyper-parameter)
batch_size = 200 # (Reduce the batch size if your computer is unable to handle
↳ it)
```

```
[7]: # Initialize the trainer class by passing the above modules
trainer = Trainer(
    dataset, optim, net, loss_func, epoch, batch_size, validate_interval=3
)
```

```
[8]: # Call the trainer function we have already implemented for you. This trains
↳ the model for the given
# hyper-parameters. It follows the same procedure as in the last ipython
↳ notebook you used for the toy-dataset
train_error, validation_accuracy = trainer.train()
```

```
Epoch Average Loss: 2.302540
Validate Acc: 0.084
Epoch Average Loss: 2.302361
Epoch Average Loss: 2.302151
Epoch Average Loss: 2.301853
Validate Acc: 0.104
Epoch Average Loss: 2.301432
Epoch Average Loss: 2.300826
Epoch Average Loss: 2.299973
Validate Acc: 0.092
Epoch Average Loss: 2.298809
Epoch Average Loss: 2.297303
Epoch Average Loss: 2.295474
Validate Acc: 0.084
Epoch Average Loss: 2.293317
Epoch Average Loss: 2.290796
Epoch Average Loss: 2.287742
Validate Acc: 0.084
Epoch Average Loss: 2.283805
Epoch Average Loss: 2.278752
Epoch Average Loss: 2.272493
Validate Acc: 0.096
Epoch Average Loss: 2.265579
Epoch Average Loss: 2.258138
Epoch Average Loss: 2.250392
```



Validate Acc: 0.108  
Epoch Average Loss: 2.242758  
Epoch Average Loss: 2.235319  
Epoch Average Loss: 2.228368  
Validate Acc: 0.112  
Epoch Average Loss: 2.221744  
Epoch Average Loss: 2.215524  
Epoch Average Loss: 2.209779  
Validate Acc: 0.124  
Epoch Average Loss: 2.204544  
Epoch Average Loss: 2.199672  
Epoch Average Loss: 2.195220  
Validate Acc: 0.136  
Epoch Average Loss: 2.191033  
Epoch Average Loss: 2.187090  
Epoch Average Loss: 2.183473  
Validate Acc: 0.140  
Epoch Average Loss: 2.179802  
Epoch Average Loss: 2.176372  
Epoch Average Loss: 2.173139  
Validate Acc: 0.140  
Epoch Average Loss: 2.170171  
Epoch Average Loss: 2.167052  
Epoch Average Loss: 2.164485  
Validate Acc: 0.140  
Epoch Average Loss: 2.161738  
Epoch Average Loss: 2.159202  
Epoch Average Loss: 2.156774  
Validate Acc: 0.144  
Epoch Average Loss: 2.154426  
Epoch Average Loss: 2.151992  
Epoch Average Loss: 2.149871  
Validate Acc: 0.148  
Epoch Average Loss: 2.148015  
Epoch Average Loss: 2.145909  
Epoch Average Loss: 2.143988  
Validate Acc: 0.148  
Epoch Average Loss: 2.142119  
Epoch Average Loss: 2.140410  
Epoch Average Loss: 2.138593  
Validate Acc: 0.148  
Epoch Average Loss: 2.136911  
Epoch Average Loss: 2.134978  
Epoch Average Loss: 2.133238  
Validate Acc: 0.152  
Epoch Average Loss: 2.132031  
Epoch Average Loss: 2.130097  
Epoch Average Loss: 2.128498

Validate Acc: 0.152  
Epoch Average Loss: 2.127047  
Epoch Average Loss: 2.125789  
Epoch Average Loss: 2.123959  
Validate Acc: 0.164  
Epoch Average Loss: 2.122555  
Epoch Average Loss: 2.120989  
Epoch Average Loss: 2.119210  
Validate Acc: 0.168  
Epoch Average Loss: 2.117758  
Epoch Average Loss: 2.116090  
Epoch Average Loss: 2.114072  
Validate Acc: 0.160  
Epoch Average Loss: 2.112825  
Epoch Average Loss: 2.111065  
Epoch Average Loss: 2.109000  
Validate Acc: 0.172  
Epoch Average Loss: 2.106939  
Epoch Average Loss: 2.104999  
Epoch Average Loss: 2.102819  
Validate Acc: 0.172  
Epoch Average Loss: 2.100567  
Epoch Average Loss: 2.098402  
Epoch Average Loss: 2.095404  
Validate Acc: 0.176  
Epoch Average Loss: 2.093283  
Epoch Average Loss: 2.090502  
Epoch Average Loss: 2.087675  
Validate Acc: 0.188  
Epoch Average Loss: 2.084531  
Epoch Average Loss: 2.081617  
Epoch Average Loss: 2.078510  
Validate Acc: 0.228  
Epoch Average Loss: 2.075273  
Epoch Average Loss: 2.072233  
Epoch Average Loss: 2.069277  
Validate Acc: 0.220  
Epoch Average Loss: 2.065468  
Epoch Average Loss: 2.062220  
Epoch Average Loss: 2.059001  
Validate Acc: 0.224  
Epoch Average Loss: 2.056238  
Epoch Average Loss: 2.053032  
Epoch Average Loss: 2.050105  
Validate Acc: 0.244  
Epoch Average Loss: 2.047188  
Epoch Average Loss: 2.044075  
Epoch Average Loss: 2.041257

Validate Acc: 0.248  
Epoch Average Loss: 2.038873  
Epoch Average Loss: 2.035811  
Epoch Average Loss: 2.032747  
Validate Acc: 0.240  
Epoch Average Loss: 2.030543  
Epoch Average Loss: 2.028426  
Epoch Average Loss: 2.025668  
Validate Acc: 0.268  
Epoch Average Loss: 2.023039  
Epoch Average Loss: 2.020892  
Epoch Average Loss: 2.019254  
Validate Acc: 0.260  
Epoch Average Loss: 2.016086  
Epoch Average Loss: 2.014286  
Epoch Average Loss: 2.011895  
Validate Acc: 0.268  
Epoch Average Loss: 2.009652  
Epoch Average Loss: 2.008635  
Epoch Average Loss: 2.006483  
Validate Acc: 0.272  
Epoch Average Loss: 2.004348  
Epoch Average Loss: 2.001680  
Epoch Average Loss: 2.000448  
Validate Acc: 0.264  
Epoch Average Loss: 1.998511  
Epoch Average Loss: 1.996869  
Epoch Average Loss: 1.994683  
Validate Acc: 0.272  
Epoch Average Loss: 1.992685  
Epoch Average Loss: 1.990515  
Epoch Average Loss: 1.989023  
Validate Acc: 0.280  
Epoch Average Loss: 1.987115  
Epoch Average Loss: 1.984757  
Epoch Average Loss: 1.983362  
Validate Acc: 0.276  
Epoch Average Loss: 1.980770  
Epoch Average Loss: 1.978400  
Epoch Average Loss: 1.976779  
Validate Acc: 0.288  
Epoch Average Loss: 1.974880  
Epoch Average Loss: 1.973159  
Epoch Average Loss: 1.970188  
Validate Acc: 0.292  
Epoch Average Loss: 1.967808  
Epoch Average Loss: 1.965204  
Epoch Average Loss: 1.962598

Validate Acc: 0.292  
Epoch Average Loss: 1.959931  
Epoch Average Loss: 1.956976  
Epoch Average Loss: 1.954877  
Validate Acc: 0.284  
Epoch Average Loss: 1.951932  
Epoch Average Loss: 1.948408  
Epoch Average Loss: 1.944597  
Validate Acc: 0.296  
Epoch Average Loss: 1.943494  
Epoch Average Loss: 1.939637  
Epoch Average Loss: 1.937788  
Validate Acc: 0.296  
Epoch Average Loss: 1.934564  
Epoch Average Loss: 1.931541  
Epoch Average Loss: 1.930091  
Validate Acc: 0.268  
Epoch Average Loss: 1.926533  
Epoch Average Loss: 1.925099  
Epoch Average Loss: 1.922507  
Validate Acc: 0.272  
Epoch Average Loss: 1.920127  
Epoch Average Loss: 1.918304  
Epoch Average Loss: 1.916498  
Validate Acc: 0.288  
Epoch Average Loss: 1.913333  
Epoch Average Loss: 1.911053  
Epoch Average Loss: 1.908856  
Validate Acc: 0.284  
Epoch Average Loss: 1.907668  
Epoch Average Loss: 1.905090  
Epoch Average Loss: 1.902684  
Validate Acc: 0.280  
Epoch Average Loss: 1.900746  
Epoch Average Loss: 1.898557  
Epoch Average Loss: 1.897128  
Validate Acc: 0.308  
Epoch Average Loss: 1.894018  
Validate Acc: 0.308  
Epoch Average Loss: 1.887973  
Epoch Average Loss: 1.887830  
Epoch Average Loss: 1.885539  
Validate Acc: 0.308  
Epoch Average Loss: 1.883132  
Epoch Average Loss: 1.879233  
Epoch Average Loss: 1.880205  
Validate Acc: 0.300  
Epoch Average Loss: 1.875765

Epoch Average Loss: 1.874388  
Epoch Average Loss: 1.872908  
Validate Acc: 0.292  
Epoch Average Loss: 1.869481  
Epoch Average Loss: 1.868074  
Epoch Average Loss: 1.867479  
Validate Acc: 0.300  
Epoch Average Loss: 1.863720  
Epoch Average Loss: 1.862010  
Epoch Average Loss: 1.860678  
Validate Acc: 0.300  
Epoch Average Loss: 1.859408  
Epoch Average Loss: 1.857180  
Epoch Average Loss: 1.854840  
Validate Acc: 0.300  
Epoch Average Loss: 1.853317  
Epoch Average Loss: 1.850804  
Epoch Average Loss: 1.849211  
Validate Acc: 0.304  
Epoch Average Loss: 1.847043  
Epoch Average Loss: 1.846299  
Epoch Average Loss: 1.842838  
Validate Acc: 0.316  
Epoch Average Loss: 1.841435  
Epoch Average Loss: 1.840141  
Epoch Average Loss: 1.837533  
Validate Acc: 0.304  
Epoch Average Loss: 1.837839  
Epoch Average Loss: 1.835047  
Epoch Average Loss: 1.832190  
Validate Acc: 0.324  
Epoch Average Loss: 1.830337  
Epoch Average Loss: 1.828371  
Epoch Average Loss: 1.826309  
Validate Acc: 0.312  
Epoch Average Loss: 1.824816  
Epoch Average Loss: 1.822740  
Epoch Average Loss: 1.822450  
Validate Acc: 0.316  
Epoch Average Loss: 1.818955  
Epoch Average Loss: 1.816644  
Epoch Average Loss: 1.815093  
Validate Acc: 0.312  
Epoch Average Loss: 1.812639  
Epoch Average Loss: 1.811698  
Epoch Average Loss: 1.810282  
Validate Acc: 0.320  
Epoch Average Loss: 1.807828

### 1.2.1 Print the training and validation accuracies for the default hyper-parameters provided

```
[9]: from ece285.utils.evaluation import get_classification_accuracy

out_train = net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)
```

```
Training acc:  0.3476
Validation acc: 0.328
```

### 1.2.2 Debug the training


With the default parameters we provided above, you should get a validation accuracy of around ~0.2 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the training loss function and the validation accuracies during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
[10]: # Plot the training loss function and validation accuracies
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()
```



```
[7]: from ece285.utils.vis_utils import visualize_grid

# Credits: http://cs231n.stanford.edu/

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net._modules[0].parameters[0]
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype("uint8"))
    plt.gca().axis("off")
    plt.show()
```

```
[11]: show_net_weights(net)
```



## 2 Tune your hyperparameters (50%)

**What's wrong?.** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength.



**Approximate results.** You should be aim to achieve a classification accuracy of greater than 40% on the validation set. Our best network gets over 40% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on cifar10 as you can (40% could serve as a reference), with a fully-connected Neural Network.

**Explain your hyperparameter tuning process below.**

**Your Answer:** I have taken three values for each parameters and run them separately.  $lr = [0.005, 0.01, 0.1]$   $weight\_decay = [0.01, 0.001, 0.005]$   $epoch = [200, 250, 300]$   $hidden\_size = [200, 220, 250]$  Out of which, I chose the hyperparameters for our model. The best parameters is for the following configuration of hyperparameters:  $lr = 0.01$   $weight\_decay = 0.005$   $epoch = 300$   $hidden\_size = 220$

I have given below a three such iterations I had done as example.

```
[6]: best_net_hyperparams = None # store the best model into this

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
  ↳#
# model hyperparams in best_net.
  ↳#
#
  ↳#
# To help debug your network, it may help to use visualizations similar to the
  ↳#
# ones we used above; these visualizations will have significant qualitative
  ↳#
# differences from the ones we saw above for the poorly tuned network.
  ↳#
#
  ↳#
# You are now free to test different combinations of hyperparameters to build
  ↳#
# various models and test them according to the above plots and visualization
  ↳#

# TODO: Show the above plots and visualizations for the default params (already
  ↳#
# done) and the best hyper-params you obtain. You only need to show this for 2
  ↳#
# sets of hyper-params.
  ↳#
# You just need to store values for the hyperparameters in best_net_hyperparams
  ↳#
```

```

# as a list in the order
# best_net_hyperparams = [lr, weight_decay, epoch, hidden_size]
#####

lr = 0.005
weight_decay = 0.01
epoch = 300
hidden_size = 250

dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
net = init_model()
optim = SGD(net, lr, weight_decay)
loss_func = CrossEntropyLoss()
batch_size = 200 # (Reduce the batch size if your computer is unable to handle
↳ it)

trainer = Trainer(
    dataset, optim, net, loss_func, epoch, batch_size, validate_interval=3
)

train_error, validation_accuracy = trainer.train()

out_train = net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)

```

```

Epoch Average Loss: 2.302526
Validate Acc: 0.100
Epoch Average Loss: 2.302349
Epoch Average Loss: 2.302157
Epoch Average Loss: 2.301955
Validate Acc: 0.100
Epoch Average Loss: 2.301731
Epoch Average Loss: 2.301477
Epoch Average Loss: 2.301186
Validate Acc: 0.096
Epoch Average Loss: 2.300845
Epoch Average Loss: 2.300439
Epoch Average Loss: 2.299985
Validate Acc: 0.100
Epoch Average Loss: 2.299462
Epoch Average Loss: 2.298887
Epoch Average Loss: 2.298226

```

Validate Acc: 0.100  
Epoch Average Loss: 2.295913  
Validate Acc: 0.096  
Epoch Average Loss: 2.295035  
Epoch Average Loss: 2.294081  
Epoch Average Loss: 2.293082  
Validate Acc: 0.100  
Epoch Average Loss: 2.291980  
Epoch Average Loss: 2.290742  
Epoch Average Loss: 2.289412  
Validate Acc: 0.084  
Epoch Average Loss: 2.287895  
Epoch Average Loss: 2.286176  
Epoch Average Loss: 2.284228  
Validate Acc: 0.084  
Epoch Average Loss: 2.281985  
Epoch Average Loss: 2.279427  
Epoch Average Loss: 2.276528  
Validate Acc: 0.096  
Epoch Average Loss: 2.273350  
Epoch Average Loss: 2.269966  
Epoch Average Loss: 2.266377  
Validate Acc: 0.096  
Epoch Average Loss: 2.262713  
Epoch Average Loss: 2.258927  
Epoch Average Loss: 2.255181  
Validate Acc: 0.096  
Epoch Average Loss: 2.251291  
Epoch Average Loss: 2.247433  
Epoch Average Loss: 2.243567  
Validate Acc: 0.108  
Epoch Average Loss: 2.239760  
Epoch Average Loss: 2.236036  
Epoch Average Loss: 2.232396  
Validate Acc: 0.112  
Epoch Average Loss: 2.228895  
Epoch Average Loss: 2.225436  
Epoch Average Loss: 2.222147  
Validate Acc: 0.120  
Epoch Average Loss: 2.218992  
Epoch Average Loss: 2.215883  
Epoch Average Loss: 2.212952  
Validate Acc: 0.124  
Epoch Average Loss: 2.210105  
Epoch Average Loss: 2.207410  
Epoch Average Loss: 2.204794  
Validate Acc: 0.124  
Epoch Average Loss: 2.202271

Epoch Average Loss: 2.199806  
Epoch Average Loss: 2.197549  
Validate Acc: 0.132  
Epoch Average Loss: 2.195255  
Epoch Average Loss: 2.193093  
Epoch Average Loss: 2.190996  
Validate Acc: 0.132  
Epoch Average Loss: 2.189070  
Epoch Average Loss: 2.187109  
Epoch Average Loss: 2.185097  
Validate Acc: 0.132  
Epoch Average Loss: 2.183253  
Epoch Average Loss: 2.181414  
Epoch Average Loss: 2.179774  
Validate Acc: 0.136  
Epoch Average Loss: 2.178070  
Epoch Average Loss: 2.176245  
Epoch Average Loss: 2.174605  
Validate Acc: 0.140  
Epoch Average Loss: 2.172976  
Epoch Average Loss: 2.171556  
Epoch Average Loss: 2.169855  
Validate Acc: 0.144  
Epoch Average Loss: 2.168502  
Epoch Average Loss: 2.167123  
Epoch Average Loss: 2.165573  
Validate Acc: 0.136  
Epoch Average Loss: 2.164206  
Epoch Average Loss: 2.162870  
Epoch Average Loss: 2.161557  
Validate Acc: 0.140  
Epoch Average Loss: 2.160264  
Epoch Average Loss: 2.159208  
Epoch Average Loss: 2.157845  
Validate Acc: 0.144  
Epoch Average Loss: 2.156593  
Epoch Average Loss: 2.155301  
Epoch Average Loss: 2.154261  
Validate Acc: 0.148  
Epoch Average Loss: 2.153046  
Epoch Average Loss: 2.151867  
Epoch Average Loss: 2.150792  
Validate Acc: 0.156  
Epoch Average Loss: 2.149738  
Epoch Average Loss: 2.148551  
Epoch Average Loss: 2.147613  
Validate Acc: 0.148  
Epoch Average Loss: 2.146707

Epoch Average Loss: 2.145632  
Epoch Average Loss: 2.144618  
Validate Acc: 0.148  
Epoch Average Loss: 2.143601  
Epoch Average Loss: 2.142692  
Epoch Average Loss: 2.141743  
Validate Acc: 0.152  
Epoch Average Loss: 2.140732  
Epoch Average Loss: 2.139929  
Epoch Average Loss: 2.138949  
Validate Acc: 0.156  
Epoch Average Loss: 2.137964  
Epoch Average Loss: 2.137243  
Epoch Average Loss: 2.136265  
Validate Acc: 0.156  
Epoch Average Loss: 2.135400  
Epoch Average Loss: 2.134447  
Epoch Average Loss: 2.133841  
Validate Acc: 0.156  
Epoch Average Loss: 2.132902  
Epoch Average Loss: 2.132003  
Epoch Average Loss: 2.131105  
Validate Acc: 0.156  
Epoch Average Loss: 2.130283  
Epoch Average Loss: 2.129488  
Epoch Average Loss: 2.128787  
Validate Acc: 0.152  
Epoch Average Loss: 2.127793  
Epoch Average Loss: 2.127045  
Epoch Average Loss: 2.126136  
Validate Acc: 0.156  
Epoch Average Loss: 2.125544  
Validate Acc: 0.164  
Epoch Average Loss: 2.122924  
Epoch Average Loss: 2.122021  
Epoch Average Loss: 2.121062  
Validate Acc: 0.164  
Epoch Average Loss: 2.120241  
Epoch Average Loss: 2.119401  
Epoch Average Loss: 2.118580  
Validate Acc: 0.168  
Epoch Average Loss: 2.117596  
Epoch Average Loss: 2.116683  
Epoch Average Loss: 2.115751  
Validate Acc: 0.176  
Epoch Average Loss: 2.114867  
Epoch Average Loss: 2.113860  
Epoch Average Loss: 2.113115

Validate Acc: 0.176  
Epoch Average Loss: 2.111935  
Epoch Average Loss: 2.110977  
Epoch Average Loss: 2.105825  
Epoch Average Loss: 2.104365  
Epoch Average Loss: 2.103172  
Validate Acc: 0.172  
Epoch Average Loss: 2.102016  
Epoch Average Loss: 2.100846  
Epoch Average Loss: 2.099578  
Validate Acc: 0.196  
Epoch Average Loss: 2.098230  
Epoch Average Loss: 2.096877  
Epoch Average Loss: 2.095460  
Validate Acc: 0.192  
Epoch Average Loss: 2.094212  
Validate Acc: 0.196  
Epoch Average Loss: 2.089738  
Epoch Average Loss: 2.088225  
Epoch Average Loss: 2.086720  
Validate Acc: 0.212  
Epoch Average Loss: 2.085139  
Epoch Average Loss: 2.083597  
Epoch Average Loss: 2.082021  
Validate Acc: 0.232  
Epoch Average Loss: 2.080342  
Epoch Average Loss: 2.078760  
Epoch Average Loss: 2.077136  
Validate Acc: 0.232  
Epoch Average Loss: 2.075617  
Epoch Average Loss: 2.073754  
Epoch Average Loss: 2.072305  
Validate Acc: 0.224  
Epoch Average Loss: 2.070566  
Epoch Average Loss: 2.068961  
Epoch Average Loss: 2.067367  
Validate Acc: 0.212  
Epoch Average Loss: 2.065681  
Epoch Average Loss: 2.063937  
Epoch Average Loss: 2.062471  
Validate Acc: 0.228  
Epoch Average Loss: 2.060685  
Epoch Average Loss: 2.059210  
Epoch Average Loss: 2.057739  
Validate Acc: 0.236  
Epoch Average Loss: 2.056144  
Epoch Average Loss: 2.054395  
Epoch Average Loss: 2.052909

Validate Acc: 0.236  
Epoch Average Loss: 2.051298  
Epoch Average Loss: 2.049977  
Epoch Average Loss: 2.048219  
Validate Acc: 0.240  
Epoch Average Loss: 2.046737  
Epoch Average Loss: 2.045331  
Epoch Average Loss: 2.043901  
Validate Acc: 0.248  
Epoch Average Loss: 2.042262  
Epoch Average Loss: 2.040729  
Epoch Average Loss: 2.039529  
Validate Acc: 0.240  
Epoch Average Loss: 2.038338  
Epoch Average Loss: 2.036717  
Epoch Average Loss: 2.035407  
Validate Acc: 0.244  
Epoch Average Loss: 2.034201  
Epoch Average Loss: 2.032762  
Epoch Average Loss: 2.031650  
Validate Acc: 0.260  
Epoch Average Loss: 2.030036  
Epoch Average Loss: 2.029005  
Epoch Average Loss: 2.027716  
Validate Acc: 0.248  
Epoch Average Loss: 2.026555  
Epoch Average Loss: 2.025246  
Epoch Average Loss: 2.023879  
Validate Acc: 0.264  
Epoch Average Loss: 2.022807  
Epoch Average Loss: 2.021365  
Epoch Average Loss: 2.020400  
Validate Acc: 0.256  
Epoch Average Loss: 2.019019  
Epoch Average Loss: 2.017685  
Epoch Average Loss: 2.016481  
Validate Acc: 0.260  
Epoch Average Loss: 2.015223  
Epoch Average Loss: 2.013901  
Epoch Average Loss: 2.012642  
Validate Acc: 0.264  
Epoch Average Loss: 2.011623  
Epoch Average Loss: 2.010306  
Epoch Average Loss: 2.008473  
Validate Acc: 0.260  
Epoch Average Loss: 2.007575  
Epoch Average Loss: 2.006137  
Epoch Average Loss: 2.004560

Validate Acc: 0.264  
Epoch Average Loss: 2.003235  
Epoch Average Loss: 2.001494  
Epoch Average Loss: 1.999827  
Validate Acc: 0.284  
Epoch Average Loss: 1.998114  
Epoch Average Loss: 1.996309  
Epoch Average Loss: 1.994329  
Validate Acc: 0.288  
Epoch Average Loss: 1.992309  
Epoch Average Loss: 1.990600  
Epoch Average Loss: 1.988072  
Validate Acc: 0.272  
Epoch Average Loss: 1.986238  
Epoch Average Loss: 1.984061  
Epoch Average Loss: 1.981524  
Validate Acc: 0.272  
Epoch Average Loss: 1.979365  
Epoch Average Loss: 1.977379  
Epoch Average Loss: 1.974764  
Validate Acc: 0.276  
Epoch Average Loss: 1.972478  
Epoch Average Loss: 1.970383  
Epoch Average Loss: 1.968435  
Validate Acc: 0.276  
Epoch Average Loss: 1.966031  
Epoch Average Loss: 1.964479  
Epoch Average Loss: 1.962515  
Validate Acc: 0.264  
Epoch Average Loss: 1.960328  
Epoch Average Loss: 1.958778  
Epoch Average Loss: 1.956944  
Validate Acc: 0.264  
Epoch Average Loss: 1.955720  
Epoch Average Loss: 1.954605  
Epoch Average Loss: 1.952892  
Validate Acc: 0.260  
Epoch Average Loss: 1.951665  
Epoch Average Loss: 1.950396  
Epoch Average Loss: 1.948771  
Validate Acc: 0.264  
Epoch Average Loss: 1.947941  
Epoch Average Loss: 1.946434  
Epoch Average Loss: 1.945457  
Validate Acc: 0.272  
Epoch Average Loss: 1.944042  
Epoch Average Loss: 1.943626  
Epoch Average Loss: 1.942291



Validate Acc: 0.276  
Epoch Average Loss: 1.940825  
Epoch Average Loss: 1.939869  
Epoch Average Loss: 1.938971  
Validate Acc: 0.280  
Epoch Average Loss: 1.937665  
Epoch Average Loss: 1.936849  
Epoch Average Loss: 1.936099  
Validate Acc: 0.272  
Epoch Average Loss: 1.935049  
Epoch Average Loss: 1.934202  
Epoch Average Loss: 1.933168  
Validate Acc: 0.276  
Epoch Average Loss: 1.932452  
Epoch Average Loss: 1.931721  
Epoch Average Loss: 1.929866  
Validate Acc: 0.296  
Epoch Average Loss: 1.929288  
Epoch Average Loss: 1.928296  
Epoch Average Loss: 1.927513  
Validate Acc: 0.288  
Epoch Average Loss: 1.926589  
Epoch Average Loss: 1.925377  
Epoch Average Loss: 1.924597  
Validate Acc: 0.284  
Epoch Average Loss: 1.924021  
Epoch Average Loss: 1.922756  
Epoch Average Loss: 1.921668  
Validate Acc: 0.288  
Epoch Average Loss: 1.920435  
Epoch Average Loss: 1.919496  
Epoch Average Loss: 1.918351  
Validate Acc: 0.292  
Epoch Average Loss: 1.917581  
Epoch Average Loss: 1.916469  
Epoch Average Loss: 1.915281  
Validate Acc: 0.296  
Epoch Average Loss: 1.914183  
Epoch Average Loss: 1.912858  
Epoch Average Loss: 1.912609  
Validate Acc: 0.292  
Epoch Average Loss: 1.911271  
Epoch Average Loss: 1.910731  
Epoch Average Loss: 1.908893  
Validate Acc: 0.296  
Epoch Average Loss: 1.908719  
Epoch Average Loss: 1.907627  
Epoch Average Loss: 1.905909

```

Validate Acc: 0.292
Epoch Average Loss: 1.905691
Epoch Average Loss: 1.903903
Epoch Average Loss: 1.902524
Validate Acc: 0.300
Epoch Average Loss: 1.902162
Epoch Average Loss: 1.900993
Epoch Average Loss: 1.899598
Validate Acc: 0.288
Epoch Average Loss: 1.899494
Epoch Average Loss: 1.897645
Epoch Average Loss: 1.896247
Validate Acc: 0.300
Epoch Average Loss: 1.895960
Epoch Average Loss: 1.894918
Epoch Average Loss: 1.893686
Validate Acc: 0.288
Epoch Average Loss: 1.892561
Epoch Average Loss: 1.891756
Epoch Average Loss: 1.890573
Validate Acc: 0.288
Epoch Average Loss: 1.888629
Epoch Average Loss: 1.888515
Epoch Average Loss: 1.886917
Validate Acc: 0.296
Epoch Average Loss: 1.886002
Epoch Average Loss: 1.884938
Epoch Average Loss: 1.884360
Validate Acc: 0.288
Epoch Average Loss: 1.883151
Epoch Average Loss: 1.881703
Epoch Average Loss: 1.880913
Validate Acc: 0.296
Epoch Average Loss: 1.879667
Epoch Average Loss: 1.878622
Training acc: 0.311
Validation acc: 0.3

```

```


[10]: from ece285.utils.vis_utils import visualize_grid
def show_net_weights(net):
    W1 = net._modules[0].parameters[0]
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype("uint8"))
    plt.gca().axis("off")
    plt.show()


```

```
[11]: # TODO: Plot the training_error and validation_accuracy of the best network (5%)
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()

# TODO: visualize the weights of the best network (5%)
show_net_weights(net)
```





```
[13]: lr = 0.01
weight_decay = 0.001
epoch = 300
hidden_size = 250

dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
net = init_model()
optim = SGD(net, lr, weight_decay)
loss_func = CrossEntropyLoss()
batch_size = 200 # (Reduce the batch size if your computer is unable to handle
↳ it)

trainer = Trainer(
```

```

    dataset, optim, net, loss_func, epoch, batch_size, validate_interval=3
)

train_error, validation_accuracy = trainer.train()

out_train = net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)

```

```

Epoch Average Loss: 2.302489
Validate Acc: 0.100
Epoch Average Loss: 2.302175
Epoch Average Loss: 2.301808
Epoch Average Loss: 2.301294
Validate Acc: 0.100
Epoch Average Loss: 2.300603
Epoch Average Loss: 2.299667
Epoch Average Loss: 2.298424
Validate Acc: 0.096
Epoch Average Loss: 2.296938
Epoch Average Loss: 2.295239
Epoch Average Loss: 2.293194
Validate Acc: 0.100
Epoch Average Loss: 2.290830
Epoch Average Loss: 2.287830
Epoch Average Loss: 2.284016
Validate Acc: 0.088
Epoch Average Loss: 2.279108
Epoch Average Loss: 2.272916
Epoch Average Loss: 2.265835
Validate Acc: 0.096
Epoch Average Loss: 2.258275
Epoch Average Loss: 2.250465
Epoch Average Loss: 2.242635
Validate Acc: 0.108
Epoch Average Loss: 2.234885
Epoch Average Loss: 2.227745
Epoch Average Loss: 2.220971
Validate Acc: 0.124
Epoch Average Loss: 2.214528
Epoch Average Loss: 2.208796
Epoch Average Loss: 2.203476
Validate Acc: 0.128

```

Epoch Average Loss: 2.198347  
Epoch Average Loss: 2.193775  
Epoch Average Loss: 2.189463  
Validate Acc: 0.132  
Epoch Average Loss: 2.185296  
Epoch Average Loss: 2.181654  
Epoch Average Loss: 2.177800  
Validate Acc: 0.140  
Epoch Average Loss: 2.174524  
Epoch Average Loss: 2.171149  
Epoch Average Loss: 2.168381  
Validate Acc: 0.144  
Epoch Average Loss: 2.165116  
Epoch Average Loss: 2.162130  
Epoch Average Loss: 2.159715  
Validate Acc: 0.140  
Epoch Average Loss: 2.157130  
Epoch Average Loss: 2.154534  
Epoch Average Loss: 2.151824  
Validate Acc: 0.148  
Epoch Average Loss: 2.149558  
Epoch Average Loss: 2.147127  
Epoch Average Loss: 2.145511  
Validate Acc: 0.152  
Epoch Average Loss: 2.142997  
Epoch Average Loss: 2.140888  
Epoch Average Loss: 2.138844  
Validate Acc: 0.148  
Epoch Average Loss: 2.137320  
Epoch Average Loss: 2.135263  
Epoch Average Loss: 2.133492  
Validate Acc: 0.156  
Epoch Average Loss: 2.131207  
Epoch Average Loss: 2.130287  
Epoch Average Loss: 2.128015  
Validate Acc: 0.164  
Epoch Average Loss: 2.126481  
Epoch Average Loss: 2.124808  
Epoch Average Loss: 2.122231  
Validate Acc: 0.144  
Epoch Average Loss: 2.121117  
Epoch Average Loss: 2.119410  
Epoch Average Loss: 2.117811  
Validate Acc: 0.156  
Epoch Average Loss: 2.115939  
Epoch Average Loss: 2.114163  
Epoch Average Loss: 2.111996  
Validate Acc: 0.168

Epoch Average Loss: 2.110252  
Epoch Average Loss: 2.107744  
Epoch Average Loss: 2.105854  
Validate Acc: 0.164  
Epoch Average Loss: 2.103644  
Epoch Average Loss: 2.101105  
Epoch Average Loss: 2.098240  
Validate Acc: 0.180  
Epoch Average Loss: 2.095810  
Epoch Average Loss: 2.092167  
Epoch Average Loss: 2.090272  
Validate Acc: 0.192  
Epoch Average Loss: 2.087279  
Epoch Average Loss: 2.083990  
Epoch Average Loss: 2.080885  
Validate Acc: 0.228  
Epoch Average Loss: 2.077561  
Epoch Average Loss: 2.074329  
Epoch Average Loss: 2.070946  
Validate Acc: 0.232  
Epoch Average Loss: 2.067553  
Epoch Average Loss: 2.063935  
Epoch Average Loss: 2.060616  
Validate Acc: 0.228  
Epoch Average Loss: 2.057186  
Epoch Average Loss: 2.053586  
Epoch Average Loss: 2.050185  
Validate Acc: 0.236  
Epoch Average Loss: 2.047432  
Epoch Average Loss: 2.044399  
Epoch Average Loss: 2.041291  
Validate Acc: 0.240  
Epoch Average Loss: 2.038023  
Epoch Average Loss: 2.034906  
Epoch Average Loss: 2.031714  
Validate Acc: 0.252  
Epoch Average Loss: 2.029569  
Epoch Average Loss: 2.026602  
Epoch Average Loss: 2.023012  
Validate Acc: 0.268  
Epoch Average Loss: 2.021186  
Epoch Average Loss: 2.019157  
Epoch Average Loss: 2.015975  
Validate Acc: 0.256  
Epoch Average Loss: 2.014066  
Epoch Average Loss: 2.011505  
Epoch Average Loss: 2.008893  
Validate Acc: 0.276

Epoch Average Loss: 2.006200  
Epoch Average Loss: 2.004285  
Epoch Average Loss: 2.002319  
Validate Acc: 0.276  
Epoch Average Loss: 2.000018  
Epoch Average Loss: 1.997167  
Epoch Average Loss: 1.994596  
Validate Acc: 0.272  
Epoch Average Loss: 1.992388  
Epoch Average Loss: 1.989815  
Epoch Average Loss: 1.987662  
Validate Acc: 0.280  
Epoch Average Loss: 1.984376  
Epoch Average Loss: 1.982268  
Epoch Average Loss: 1.978841  
Validate Acc: 0.280  
Epoch Average Loss: 1.976054  
Epoch Average Loss: 1.973328  
Epoch Average Loss: 1.969381  
Validate Acc: 0.288  
Epoch Average Loss: 1.964949  
Epoch Average Loss: 1.961530  
Epoch Average Loss: 1.958092  
Validate Acc: 0.268  
Epoch Average Loss: 1.954897  
Epoch Average Loss: 1.949979  
Epoch Average Loss: 1.947427  
Validate Acc: 0.276  
Epoch Average Loss: 1.943943  
Epoch Average Loss: 1.941818  
Epoch Average Loss: 1.938477  
Validate Acc: 0.268  
Epoch Average Loss: 1.936221  
Epoch Average Loss: 1.933333  
Epoch Average Loss: 1.930949  
Validate Acc: 0.272  
Epoch Average Loss: 1.928596  
Epoch Average Loss: 1.926733  
Epoch Average Loss: 1.923987  
Validate Acc: 0.276  
Epoch Average Loss: 1.921392  
Epoch Average Loss: 1.920596  
Epoch Average Loss: 1.916185  
Validate Acc: 0.292  
Epoch Average Loss: 1.915205  
Epoch Average Loss: 1.913602  
Epoch Average Loss: 1.910021  
Validate Acc: 0.276



Epoch Average Loss: 1.907992  
Epoch Average Loss: 1.906756  
Epoch Average Loss: 1.904837  
Validate Acc: 0.296  
Epoch Average Loss: 1.900463  
Epoch Average Loss: 1.899878  
Epoch Average Loss: 1.897696  
Validate Acc: 0.288  
Epoch Average Loss: 1.894868  
Epoch Average Loss: 1.892308  
Epoch Average Loss: 1.890190  
Validate Acc: 0.280  
Epoch Average Loss: 1.888111  
Epoch Average Loss: 1.885584  
Epoch Average Loss: 1.882782  
Validate Acc: 0.280  
Epoch Average Loss: 1.880600  
Epoch Average Loss: 1.880736  
Epoch Average Loss: 1.878126  
Validate Acc: 0.312  
Epoch Average Loss: 1.873254  
Epoch Average Loss: 1.871680  
Epoch Average Loss: 1.870764  
Validate Acc: 0.288  
Epoch Average Loss: 1.867330  
Epoch Average Loss: 1.865461  
Epoch Average Loss: 1.863002  
Validate Acc: 0.296  
Epoch Average Loss: 1.861743  
Epoch Average Loss: 1.859313  
Epoch Average Loss: 1.857925  
Validate Acc: 0.316  
Epoch Average Loss: 1.855220  
Epoch Average Loss: 1.852505  
Epoch Average Loss: 1.850832  
Validate Acc: 0.300  
Epoch Average Loss: 1.849007  
Epoch Average Loss: 1.846634  
Epoch Average Loss: 1.843804  
Validate Acc: 0.308  
Epoch Average Loss: 1.842197  
Validate Acc: 0.296  
Epoch Average Loss: 1.835131  
Epoch Average Loss: 1.832353  
Epoch Average Loss: 1.829254  
Validate Acc: 0.288  
Epoch Average Loss: 1.827850  
Epoch Average Loss: 1.824839

Epoch Average Loss: 1.823416  
Validate Acc: 0.312  
Epoch Average Loss: 1.821212  
Epoch Average Loss: 1.818873  
Epoch Average Loss: 1.817688  
Validate Acc: 0.312  
Epoch Average Loss: 1.813504  
Epoch Average Loss: 1.813492  
Epoch Average Loss: 1.810200  
Validate Acc: 0.316  
Epoch Average Loss: 1.808298  
Epoch Average Loss: 1.808731  
Epoch Average Loss: 1.803758  
Validate Acc: 0.300  
Epoch Average Loss: 1.802092  
Epoch Average Loss: 1.800509  
Epoch Average Loss: 1.798266  
Validate Acc: 0.316  
Epoch Average Loss: 1.794417  
Epoch Average Loss: 1.792550  
Epoch Average Loss: 1.789978  
Validate Acc: 0.332  
Epoch Average Loss: 1.789146  
Epoch Average Loss: 1.787217  
Epoch Average Loss: 1.784517  
Validate Acc: 0.320  
Epoch Average Loss: 1.781417  
Epoch Average Loss: 1.778753  
Epoch Average Loss: 1.776784  
Validate Acc: 0.344  
Epoch Average Loss: 1.774150  
Epoch Average Loss: 1.773725  
Epoch Average Loss: 1.774651  
Validate Acc: 0.320  
Epoch Average Loss: 1.773029  
Epoch Average Loss: 1.768910  
Epoch Average Loss: 1.766151  
Validate Acc: 0.324  
Epoch Average Loss: 1.764716  
Epoch Average Loss: 1.762030  
Epoch Average Loss: 1.760418  
Validate Acc: 0.328  
Epoch Average Loss: 1.755896  
Epoch Average Loss: 1.757062  
Epoch Average Loss: 1.754094  
Validate Acc: 0.360  
Epoch Average Loss: 1.752389  
Epoch Average Loss: 1.750901

Epoch Average Loss: 1.747780  
Validate Acc: 0.336  
Epoch Average Loss: 1.746178  
Epoch Average Loss: 1.743281  
Epoch Average Loss: 1.742734  
Validate Acc: 0.348  
Epoch Average Loss: 1.739866  
Epoch Average Loss: 1.737988  
Epoch Average Loss: 1.738288  
Validate Acc: 0.360  
Epoch Average Loss: 1.735300  
Epoch Average Loss: 1.731379  
Epoch Average Loss: 1.733487  
Validate Acc: 0.360  
Epoch Average Loss: 1.729472  
Epoch Average Loss: 1.727046  
Epoch Average Loss: 1.725056  
Validate Acc: 0.360  
Epoch Average Loss: 1.723247  
Epoch Average Loss: 1.722433  
Epoch Average Loss: 1.721847  
Validate Acc: 0.368  
Epoch Average Loss: 1.717599  
Epoch Average Loss: 1.717074  
Epoch Average Loss: 1.716074  
Validate Acc: 0.384  
Epoch Average Loss: 1.711912  
Epoch Average Loss: 1.710651  
Epoch Average Loss: 1.707093  
Validate Acc: 0.364  
Epoch Average Loss: 1.707324  
Epoch Average Loss: 1.705904  
Epoch Average Loss: 1.704887  
Validate Acc: 0.372  
Epoch Average Loss: 1.704297  
Epoch Average Loss: 1.701699  
Epoch Average Loss: 1.698358  
Validate Acc: 0.364  
Epoch Average Loss: 1.697639  
Epoch Average Loss: 1.698685  
Epoch Average Loss: 1.692287  
Validate Acc: 0.372  
Epoch Average Loss: 1.692875  
Epoch Average Loss: 1.688885  
Epoch Average Loss: 1.688290  
Validate Acc: 0.380  
Epoch Average Loss: 1.686823  
Epoch Average Loss: 1.686121


Epoch Average Loss: 1.684055  
Validate Acc: 0.388  
Epoch Average Loss: 1.682240  
Epoch Average Loss: 1.677842  
Epoch Average Loss: 1.677588  
Validate Acc: 0.388  
Epoch Average Loss: 1.677669  
Epoch Average Loss: 1.676395  
Epoch Average Loss: 1.673974  
Validate Acc: 0.388  
Epoch Average Loss: 1.674011  
Epoch Average Loss: 1.675739  
Epoch Average Loss: 1.670373  
Validate Acc: 0.392  
Epoch Average Loss: 1.667458  
Epoch Average Loss: 1.669270  
Epoch Average Loss: 1.664216  
Validate Acc: 0.396  
Epoch Average Loss: 1.662761  
Epoch Average Loss: 1.660724  
Epoch Average Loss: 1.661186  
Validate Acc: 0.408  
Epoch Average Loss: 1.659134  
Epoch Average Loss: 1.658604  
Epoch Average Loss: 1.653877  
Validate Acc: 0.404  
Epoch Average Loss: 1.655886  
Epoch Average Loss: 1.650004  
Epoch Average Loss: 1.653319  
Validate Acc: 0.408  
Epoch Average Loss: 1.649841  
Epoch Average Loss: 1.646097  
Epoch Average Loss: 1.646471  
Validate Acc: 0.396  
Epoch Average Loss: 1.644635  
Epoch Average Loss: 1.641709  
Epoch Average Loss: 1.640410  
Validate Acc: 0.404  
Epoch Average Loss: 1.637459  
Epoch Average Loss: 1.639823  
Epoch Average Loss: 1.636308  
Validate Acc: 0.400  
Epoch Average Loss: 1.631699  
Epoch Average Loss: 1.634376  
Epoch Average Loss: 1.632244  
Validate Acc: 0.420  
Epoch Average Loss: 1.632382  
Epoch Average Loss: 1.630535


```
Epoch Average Loss: 1.627151
Validate Acc: 0.412
Epoch Average Loss: 1.629897
Epoch Average Loss: 1.628444
Epoch Average Loss: 1.623241
Validate Acc: 0.408
Epoch Average Loss: 1.621192
Epoch Average Loss: 1.617033
Epoch Average Loss: 1.617744
Validate Acc: 0.424
Epoch Average Loss: 1.617058
Epoch Average Loss: 1.615454
Epoch Average Loss: 1.615704
Validate Acc: 0.400
Epoch Average Loss: 1.614581
Epoch Average Loss: 1.615586
Epoch Average Loss: 1.611197
Validate Acc: 0.376
Epoch Average Loss: 1.606902
Epoch Average Loss: 1.605216
Epoch Average Loss: 1.604763
Validate Acc: 0.420
Epoch Average Loss: 1.603948
Epoch Average Loss: 1.605291
Epoch Average Loss: 1.602218
Validate Acc: 0.424
Epoch Average Loss: 1.599011
Epoch Average Loss: 1.597734
Training acc: 0.4318
Validation acc: 0.428
```

```
[14]: # TODO: Plot the training_error and validation_accuracy of the best network (5%)
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()
```

```
# TODO: visualize the weights of the best network (5%)
show_net_weights(net)
```





```
[12]: lr = 0.01
weight_decay = 0.005
epoch = 300
hidden_size = 220

dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
best_net = init_model()
optim = SGD(best_net, lr, weight_decay)
loss_func = CrossEntropyLoss()
batch_size = 200 # (Reduce the batch size if your computer is unable to handle
→ it)
```

```

trainer = Trainer(
    dataset, optim, best_net, loss_func, epoch, batch_size, validate_interval=3
)

train_error, validation_accuracy = trainer.train()

out_train = best_net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = best_net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)

```

```

Epoch Average Loss: 2.302472
Validate Acc: 0.088
Epoch Average Loss: 2.302144
Epoch Average Loss: 2.301731
Epoch Average Loss: 2.301198
Validate Acc: 0.096
Epoch Average Loss: 2.300464
Epoch Average Loss: 2.299494
Epoch Average Loss: 2.298254
Validate Acc: 0.092
Epoch Average Loss: 2.296750
Epoch Average Loss: 2.294986
Epoch Average Loss: 2.292963
Validate Acc: 0.084
Epoch Average Loss: 2.290522
Epoch Average Loss: 2.287571
Epoch Average Loss: 2.283761
Validate Acc: 0.088
Epoch Average Loss: 2.278771
Epoch Average Loss: 2.272477
Epoch Average Loss: 2.265415
Validate Acc: 0.096
Epoch Average Loss: 2.257865
Epoch Average Loss: 2.250176
Epoch Average Loss: 2.242396
Validate Acc: 0.108
Epoch Average Loss: 2.234918
Epoch Average Loss: 2.227660
Epoch Average Loss: 2.221030
Validate Acc: 0.120
Epoch Average Loss: 2.214699
Epoch Average Loss: 2.209097
Epoch Average Loss: 2.203624

```



Validate Acc: 0.124  
Epoch Average Loss: 2.199013  
Epoch Average Loss: 2.194167  
Epoch Average Loss: 2.189933  
Validate Acc: 0.132  
Epoch Average Loss: 2.185788  
Epoch Average Loss: 2.182165  
Epoch Average Loss: 2.178585  
Validate Acc: 0.136  
Epoch Average Loss: 2.175092  
Epoch Average Loss: 2.171694  
Epoch Average Loss: 2.168471  
Validate Acc: 0.136  
Epoch Average Loss: 2.165718  
Epoch Average Loss: 2.163014  
Epoch Average Loss: 2.160226  
Validate Acc: 0.140  
Epoch Average Loss: 2.157685  
Epoch Average Loss: 2.155111  
Epoch Average Loss: 2.152674  
Validate Acc: 0.140  
Epoch Average Loss: 2.150621  
Epoch Average Loss: 2.148555  
Epoch Average Loss: 2.146032  
Validate Acc: 0.144  
Epoch Average Loss: 2.144091  
Epoch Average Loss: 2.141980  
Epoch Average Loss: 2.140193  
Validate Acc: 0.160  
Epoch Average Loss: 2.138349  
Epoch Average Loss: 2.136543  
Epoch Average Loss: 2.134345  
Validate Acc: 0.148  
Epoch Average Loss: 2.132810  
Epoch Average Loss: 2.131095  
Epoch Average Loss: 2.129480  
Validate Acc: 0.160  
Epoch Average Loss: 2.127636  
Epoch Average Loss: 2.126270  
Epoch Average Loss: 2.124454  
Validate Acc: 0.164  
Epoch Average Loss: 2.122775  
Epoch Average Loss: 2.121050  
Epoch Average Loss: 2.119146  
Validate Acc: 0.172  
Epoch Average Loss: 2.117335  
Epoch Average Loss: 2.115456  
Epoch Average Loss: 2.113508

Validate Acc: 0.172  
Epoch Average Loss: 2.111606  
Epoch Average Loss: 2.109921  
Epoch Average Loss: 2.107675  
Validate Acc: 0.172  
Epoch Average Loss: 2.105146  
Epoch Average Loss: 2.102636  
Epoch Average Loss: 2.100345  
Validate Acc: 0.172  
Epoch Average Loss: 2.097641  
Epoch Average Loss: 2.095070  
Epoch Average Loss: 2.092407  
Validate Acc: 0.180  
Epoch Average Loss: 2.089245  
Epoch Average Loss: 2.085959  
Epoch Average Loss: 2.082830  
Validate Acc: 0.220  
Epoch Average Loss: 2.079610  
Epoch Average Loss: 2.076454  
Epoch Average Loss: 2.073281  
Validate Acc: 0.216  
Epoch Average Loss: 2.069660  
Epoch Average Loss: 2.066625  
Epoch Average Loss: 2.062764  
Validate Acc: 0.220  
Epoch Average Loss: 2.059428  
Epoch Average Loss: 2.056640  
Epoch Average Loss: 2.053280  
Validate Acc: 0.232  
Epoch Average Loss: 2.049563  
Epoch Average Loss: 2.046491  
Epoch Average Loss: 2.043952  
Validate Acc: 0.240  
Epoch Average Loss: 2.040936  
Epoch Average Loss: 2.037675  
Epoch Average Loss: 2.035606  
Validate Acc: 0.260  
Epoch Average Loss: 2.032410  
Epoch Average Loss: 2.030140  
Epoch Average Loss: 2.027801  
Validate Acc: 0.252  
Epoch Average Loss: 2.024898  
Epoch Average Loss: 2.022218  
Epoch Average Loss: 2.019796  
Validate Acc: 0.244  
Epoch Average Loss: 2.017924  
Epoch Average Loss: 2.015195  
Epoch Average Loss: 2.012998

Validate Acc: 0.260  
Epoch Average Loss: 2.011441  
Epoch Average Loss: 2.008841  
Epoch Average Loss: 2.006403  
Validate Acc: 0.256  
Epoch Average Loss: 2.004663  
Epoch Average Loss: 2.002702  
Epoch Average Loss: 2.000946  
Validate Acc: 0.272  
Epoch Average Loss: 1.998283  
Epoch Average Loss: 1.996196  
Epoch Average Loss: 1.994684  
Validate Acc: 0.276  
Epoch Average Loss: 1.992475  
Epoch Average Loss: 1.990613  
Epoch Average Loss: 1.988349  
Validate Acc: 0.276  
Epoch Average Loss: 1.986744  
Epoch Average Loss: 1.984256  
Epoch Average Loss: 1.981858  
Validate Acc: 0.276  
Epoch Average Loss: 1.980328  
Epoch Average Loss: 1.979070  
Epoch Average Loss: 1.975789  
Validate Acc: 0.292  
Epoch Average Loss: 1.973444  
Epoch Average Loss: 1.970607  
Epoch Average Loss: 1.967387  
Validate Acc: 0.288  
Epoch Average Loss: 1.964749  
Epoch Average Loss: 1.961034  
Epoch Average Loss: 1.957519  
Validate Acc: 0.296  
Epoch Average Loss: 1.954154  
Epoch Average Loss: 1.950538  
Epoch Average Loss: 1.947467  
Validate Acc: 0.284  
Epoch Average Loss: 1.944108  
Epoch Average Loss: 1.940550  
Epoch Average Loss: 1.937393  
Validate Acc: 0.264  
Epoch Average Loss: 1.935336  
Epoch Average Loss: 1.932406  
Epoch Average Loss: 1.928246  
Validate Acc: 0.284  
Epoch Average Loss: 1.927339  
Epoch Average Loss: 1.923892  
Epoch Average Loss: 1.921886

Validate Acc: 0.292  
Epoch Average Loss: 1.917909  
Epoch Average Loss: 1.916347  
Epoch Average Loss: 1.913984  
Validate Acc: 0.284  
Epoch Average Loss: 1.910486  
Epoch Average Loss: 1.910019  
Epoch Average Loss: 1.907236  
Validate Acc: 0.288  
Epoch Average Loss: 1.906062  
Epoch Average Loss: 1.903019  
Epoch Average Loss: 1.901872  
Validate Acc: 0.292  
Epoch Average Loss: 1.899603  
Epoch Average Loss: 1.896029  
Epoch Average Loss: 1.893644  
Validate Acc: 0.288  
Epoch Average Loss: 1.891181  
Epoch Average Loss: 1.889215  
Epoch Average Loss: 1.888081  
Validate Acc: 0.308  
Epoch Average Loss: 1.884229  
Epoch Average Loss: 1.882331  
Epoch Average Loss: 1.879765  
Validate Acc: 0.312  
Epoch Average Loss: 1.878998  
Epoch Average Loss: 1.875771  
Epoch Average Loss: 1.874489  
Validate Acc: 0.288  
Epoch Average Loss: 1.872152  
Epoch Average Loss: 1.870157  
Epoch Average Loss: 1.868175  
Validate Acc: 0.304  
Epoch Average Loss: 1.865281  
Epoch Average Loss: 1.862937  
Epoch Average Loss: 1.860678  
Validate Acc: 0.296  
Epoch Average Loss: 1.857830  
Epoch Average Loss: 1.856804  
Epoch Average Loss: 1.853794  
Validate Acc: 0.288  
Epoch Average Loss: 1.852897  
Epoch Average Loss: 1.849328  
Epoch Average Loss: 1.847137  
Validate Acc: 0.324  
Epoch Average Loss: 1.847143  
Epoch Average Loss: 1.842732  
Epoch Average Loss: 1.840457

Validate Acc: 0.316  
Epoch Average Loss: 1.838717  
Epoch Average Loss: 1.837108  
Epoch Average Loss: 1.834021  
Validate Acc: 0.316  
Epoch Average Loss: 1.832217  
Epoch Average Loss: 1.830930  
Epoch Average Loss: 1.828427  
Validate Acc: 0.304  
Epoch Average Loss: 1.826136  
Epoch Average Loss: 1.822719  
Epoch Average Loss: 1.821312  
Validate Acc: 0.324  
Epoch Average Loss: 1.818760  
Epoch Average Loss: 1.815600  
Epoch Average Loss: 1.814089  
Validate Acc: 0.312  
Epoch Average Loss: 1.811374  
Epoch Average Loss: 1.808949  
Epoch Average Loss: 1.806689  
Validate Acc: 0.332  
Epoch Average Loss: 1.806491  
Epoch Average Loss: 1.802886  
Epoch Average Loss: 1.803246  
Validate Acc: 0.308  
Epoch Average Loss: 1.797665  
Epoch Average Loss: 1.796849  
Epoch Average Loss: 1.795012  
Validate Acc: 0.320  
Epoch Average Loss: 1.792252  
Epoch Average Loss: 1.790518  
Epoch Average Loss: 1.789773  
Validate Acc: 0.328  
Epoch Average Loss: 1.786349  
Epoch Average Loss: 1.783976  
Epoch Average Loss: 1.783662  
Validate Acc: 0.336  
Epoch Average Loss: 1.778963  
Epoch Average Loss: 1.778777  
Epoch Average Loss: 1.776502  
Validate Acc: 0.328  
Epoch Average Loss: 1.773648  
Epoch Average Loss: 1.771008  
Epoch Average Loss: 1.769815  
Validate Acc: 0.324  
Epoch Average Loss: 1.768022  
Epoch Average Loss: 1.765407  
Epoch Average Loss: 1.764688

Validate Acc: 0.344  
Epoch Average Loss: 1.761039  
Epoch Average Loss: 1.761452  
Epoch Average Loss: 1.758851  
Validate Acc: 0.332  
Epoch Average Loss: 1.755237  
Epoch Average Loss: 1.754111  
Epoch Average Loss: 1.752083  
Validate Acc: 0.356  
Epoch Average Loss: 1.751765  
Epoch Average Loss: 1.746875  
Epoch Average Loss: 1.747984  
Validate Acc: 0.360  
Epoch Average Loss: 1.743258  
Epoch Average Loss: 1.744117  
Epoch Average Loss: 1.742469  
Validate Acc: 0.372  
Epoch Average Loss: 1.740599  
Epoch Average Loss: 1.736201  
Epoch Average Loss: 1.735823  
Validate Acc: 0.360  
Epoch Average Loss: 1.733081  
Epoch Average Loss: 1.731261  
Epoch Average Loss: 1.730866  
Validate Acc: 0.360  
Epoch Average Loss: 1.730214  
Epoch Average Loss: 1.724968  
Epoch Average Loss: 1.725022  
Validate Acc: 0.356  
Epoch Average Loss: 1.723810  
Epoch Average Loss: 1.721130  
Epoch Average Loss: 1.719231  
Validate Acc: 0.368  
Epoch Average Loss: 1.716712  
Epoch Average Loss: 1.720263  
Epoch Average Loss: 1.714299  
Validate Acc: 0.360  
Epoch Average Loss: 1.712697  
Epoch Average Loss: 1.711832  
Epoch Average Loss: 1.711066  
Validate Acc: 0.368  
Epoch Average Loss: 1.706354  
Epoch Average Loss: 1.706054  
Epoch Average Loss: 1.706269  
Validate Acc: 0.364  
Epoch Average Loss: 1.701834  
Epoch Average Loss: 1.701507  
Epoch Average Loss: 1.702851

Validate Acc: 0.372  
Epoch Average Loss: 1.697556  
Epoch Average Loss: 1.697616  
Epoch Average Loss: 1.693791  
Validate Acc: 0.388  
Epoch Average Loss: 1.694119  
Epoch Average Loss: 1.691758  
Epoch Average Loss: 1.688598  
Validate Acc: 0.388  
Epoch Average Loss: 1.686839  
Epoch Average Loss: 1.688522  
Epoch Average Loss: 1.687408  
Validate Acc: 0.384  
Epoch Average Loss: 1.683037  
Epoch Average Loss: 1.681893  
Epoch Average Loss: 1.679489  
Validate Acc: 0.376  
Epoch Average Loss: 1.679530  
Epoch Average Loss: 1.678412  
Epoch Average Loss: 1.676471  
Validate Acc: 0.396  
Epoch Average Loss: 1.677150  
Epoch Average Loss: 1.673772  
Epoch Average Loss: 1.673934  
Validate Acc: 0.392  
Epoch Average Loss: 1.670671  
Epoch Average Loss: 1.672209  
Epoch Average Loss: 1.666345  
Validate Acc: 0.380  
Epoch Average Loss: 1.667167  
Epoch Average Loss: 1.666193  
Epoch Average Loss: 1.662227  
Validate Acc: 0.404  
Epoch Average Loss: 1.661961  
Epoch Average Loss: 1.660465  
Epoch Average Loss: 1.657293  
Validate Acc: 0.380  
Epoch Average Loss: 1.658638  
Epoch Average Loss: 1.657067  
Epoch Average Loss: 1.653397  
Validate Acc: 0.392  
Epoch Average Loss: 1.653204  
Epoch Average Loss: 1.654069  
Epoch Average Loss: 1.650875  
Validate Acc: 0.388  
Epoch Average Loss: 1.649861  
Epoch Average Loss: 1.648341  
Epoch Average Loss: 1.644715

```
Validate Acc: 0.384
Epoch Average Loss: 1.647571
Epoch Average Loss: 1.647479
Epoch Average Loss: 1.643969
Validate Acc: 0.400
Epoch Average Loss: 1.638793
Epoch Average Loss: 1.638986
Epoch Average Loss: 1.639337
Validate Acc: 0.412
Epoch Average Loss: 1.637587
Epoch Average Loss: 1.635226
Epoch Average Loss: 1.635178
Validate Acc: 0.412
Epoch Average Loss: 1.634137
Epoch Average Loss: 1.634714
Epoch Average Loss: 1.632701
Validate Acc: 0.396
Epoch Average Loss: 1.630597
Epoch Average Loss: 1.628989
Epoch Average Loss: 1.627605
Validate Acc: 0.404
Epoch Average Loss: 1.625635
Epoch Average Loss: 1.624906
Epoch Average Loss: 1.622428
Validate Acc: 0.396
Epoch Average Loss: 1.622633
Epoch Average Loss: 1.619554
Epoch Average Loss: 1.620196
Validate Acc: 0.408
Epoch Average Loss: 1.618700
Epoch Average Loss: 1.616955
Training acc: 0.427
Validation acc: 0.432
```


```
[13]: # TODO: Plot the training_error and validation_accuracy of the best network (5%)
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")


plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
```



```
plt.legend()
plt.show()

# TODO: visualize the weights of the best network (5%)
show_net_weights(best_net)
```





```
[14]: acc = get_classification_accuracy(out_train, y_train)
      print("Training acc: ", acc)
      out_val = best_net.predict(x_val)
      acc = get_classification_accuracy(out_val, y_val)
      print("Validation acc: ", acc)
```

```
Training acc:  0.427
Validation acc: 0.432
```

### 3 Run on the test set (30%)

When you are done experimenting, you should evaluate your final trained network on the test set; you should get above 35%.

```
[15]: test_acc = (best_net.predict(x_test) == y_test).mean()  
      print("Test accuracy: ", test_acc)
```

Test accuracy: 0.364

**Inline Question (10%)** Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

**Your Answer:**

1. Train on a larger dataset.
2. Increase the regularization strength.

**Your Explanation:** Expanding the dataset for training purposes can assist in narrowing the gap between training and testing accuracy, as it provides the model with a more diverse range of examples to learn from, allowing for better generalization to new and unseen data. Similarly, augmenting the regularization strength can help decrease overfitting, thus resulting in a smaller disparity between the model's training and testing accuracy.